Faculty of Informatics and Computer Science

Artificial intelligence

Integrative Real-Time Sign Language Recognition for Enhanced Communication

By: Ezzeldein Yasser Hassan

Supervised By

Dr. Nermin Othman

June 2024

1 Abstract

For many years, people who are Deaf and mute have used sign language, one of the most appreciated and ancient forms of communication, as their primary means of interaction. Sign language has played a crucial role in bridging communication gaps within the Deaf community and beyond. However, despite its significance, a large percentage of the global population remains unfamiliar with sign language, creating substantial barriers to communication with the Deaf community. This lack of familiarity often results in isolation and difficulties for Deaf individuals in various social, educational, and professional settings.

To address this pressing issue, modern machine learning and deep learning techniques have been harnessed to develop a live, real-time fingerspelling application. This innovative solution aims to facilitate communication by translating sign language into text, thereby enhancing interaction between Deaf individuals and those who do not know sign language. At the heart of this technology are gestures, which serve as the physical representations and markers of emotions within the Deaf community. These gestures are not only crucial for conveying meaning but also for expressing emotions, making them an integral part of the communication process.

The development of this application is particularly significant in our complex environment, where people with different needs, including the Deaf community, coexist with those without impairments. Sign language, recognized as one of the world's most essential languages, serves as a vital tool for bridging communication gaps among diverse groups. The importance of this language is underscored by statistics from the World Health Organization (WHO), which reported that the global population of individuals with hearing impairments numbered 278 million in 2005, a figure that soared to 466 million by 2019. This dramatic increase highlights the growing need for effective communication solutions for the Deaf community.

Our project leverages supervised machine learning/deep learning techniques to translate sign language into text, addressing the serious communication challenges faced by the Deaf community. This involves the development of a thoroughly trained model, which undergoes rigorous accuracy testing to ensure its effectiveness. The integration of data capture techniques enables the application to accept user input, further enhancing its functionality and user experience.

By bridging the communication gap through advanced technology, this project aims to foster a more inclusive society where Deaf individuals can communicate seamlessly with the wider community. The implementation of this real-time fingerspelling application represents a significant step forward in improving accessibility and interaction for the Deaf community, ultimately contributing to greater social integration and equality.

2 Turnitin Report

ezz(206879)_GP_thisis

ORIGINALITY REPORT

18%	9%	10%	12%
SIMILARITY INDEX	INTERNET SOURCES	PUBLICATIONS	STUDENT PAPERS

PRIMARY SOURCES

1	Abdelaziz Testas. "Distributed Machine Learning with PySpark", Springer Science and Business Media LLC, 2023 Publication	1%
2	Submitted to Liverpool John Moores University Student Paper	1%
3	Submitted to British University in Egypt Student Paper	1%
4	Submitted to University of Queensland Student Paper	1%
5	www.geeksforgeeks.org Internet Source	<1%
6	Submitted to Indian School of Business Student Paper	<1%
7	medium.com Internet Source	<1%
8	Submitted to Manchester Metropolitan University Student Paper	<1%

9	Submitted to University of Greenwich Student Paper	<1 %
10	Submitted to University of Northampton Student Paper	<1 %
11	Abdelaziz Testas. "Chapter 10 Support Vector Machine Classification with Pandas, Scikit-Learn, and PySpark", Springer Science and Business Media LLC, 2023 Publication	<1 %
12	www.ijisae.org Internet Source	<1 %
13	icente.selcuk.edu.tr Internet Source	<1 %
14	arxiv.org Internet Source	<1 %
15	Singh, Shekhar. "Facial Expression Recognition Using Convolutional Neural Networks (CNNs) and Generative Adversarial Networks (GANs) for Data Augmentation and Image Generation", University of Nevada, Las Vegas, 2024 Publication	<1 %
16	Eklas Hossain. "Machine Learning Crash Course for Engineers", Springer Science and Business Media LLC, 2024 Publication	<1 %

17	I.A. Adeyanju, O.O. Bello, M.A. Adegboye. "Machine learning methods for sign language recognition: A critical review and analysis", Intelligent Systems with Applications, 2021 Publication	<1 %
18	Submitted to City University of Hong Kong Student Paper	<1 %
19	Submitted to University of Wales Institute, Cardiff Student Paper	<1 %
20	www.ijirset.com Internet Source	<1 %
21	www.mdpi.com Internet Source	<1 %
22	Submitted to University of North Texas Student Paper	<1 %
23	dokumen.pub Internet Source	<1 %
24	Submitted to Intercollege Student Paper	<1 %
25	fastercapital.com Internet Source	<1 %
26	Submitted to California State University, Fresno Student Paper	<1 %

27	portalcris.vdu.lt Internet Source	<1 %
28	Submitted to CSU, Chico Student Paper	<1 %
29	Submitted to Infile Student Paper	<1 %
30	Submitted to Angeles University Foundation Student Paper	<1 %
31	Submitted to Aston University Student Paper	<1 %
32	ijereee.com Internet Source	<1 %
33	Submitted to University of Bristol Student Paper	<1 %
34	github.com Internet Source	<1 %
35	lutpub.lut.fi Internet Source	<1 %
36	Submitted to Southampton Solent University Student Paper	<1 %
37	Submitted to Asia Pacific University College of Technology and Innovation (UCTI) Student Paper	<1 %
38	Submitted to Cerritos College	

	Student Paper	<1 %
39	Submitted to Colorado Technical University Student Paper	<1 %
40	Maliha Tasnim, Tasnia Islam Rifa, Tanvir Shahriar, Mohammad Ahsan Habib. "Wind energy deployment in Bangladesh: Investigating feasible locations and their characteristics", Energy Reports, 2024 Publication	<1 %
41	Submitted to The Robert Gordon University Student Paper	<1 %
42	Submitted to University of East London Student Paper	<1 %
43	Submitted to Monash University Student Paper	<1 %
44	Submitted to South Bank University Student Paper	<1 %
45	Submitted to University of Sunderland Student Paper	<1 %
46	www.mlq.ai Internet Source	<1 %
47	www.nature.com Internet Source	<1 %
	Submitted to Associatie K.U.Leuven	

48	Student Paper	<1 %
49	Submitted to University of Glasgow Student Paper	<1 %
50	training.galaxyproject.org Internet Source	<1 %
51	www.eajournals.org Internet Source	<1 %
52	Submitted to AlHussein Technical University Student Paper	<1 %
53	www.frontiersin.org Internet Source	<1 %
54	Caio Cesar Teodoro Mendes, Vincent Fremont, Denis Fernando Wolf. "Exploiting fully convolutional neural networks for fast road detection", 2016 IEEE International Conference on Robotics and Automation (ICRA), 2016 Publication	<1 %
55	Poornachandra Sarang. "Artificial Neural Networks with TensorFlow 2", Springer Science and Business Media LLC, 2021 Publication	<1 %
56	Submitted to Taylor's Education Group Student Paper	<1 %

57	Submitted to University of Lancaster Student Paper	<1 %
58	www.ijritcc.org Internet Source	<1 %
59	Avgoustinos Filippoupolitis, Babak Takand, George Loukas. "Activity Recognition in a Home Setting Using Off the Shelf Smart Watch Technology", 2016 15th International Conference on Ubiquitous Computing and Communications and 2016 International Symposium on Cyberspace and Security (IUCC-CSS), 2016 Publication	<1 %
60	Submitted to ESC Rennes Student Paper	<1 %
61	Submitted to Federal University of Technology Student Paper	<1 %
62	Submitted to CSU Northridge Student Paper	<1 %
63	Submitted to Cranfield University Student Paper	<1 %
64	Lecture Notes in Computer Science, 2015. Publication	<1 %
65	Submitted to University of Bradford Student Paper	<1 %

66	D.G. Kodagoda, K.G.R.U Ishara, R.M.R.P Kumara, W.A.D.T Dilshan, Lakesha Weerasinghe, Nadeesa Premadasa. "An Interactive E-Learning Tool", 2022 IEEE 7th International conference for Convergence in Technology (I2CT), 2022 Publication	<1 %
67	Submitted to National College of Ireland Student Paper	<1 %
68	Submitted to University of Central Lancashire Student Paper	<1 %
69	Submitted to University of Sydney Student Paper	<1 %
70	Arunava Mukhopadhyay, Aritra Chakrabarty, Agnish Arpan Das, Aishik Sarkar. "Hand Gesture Based Recognition System", 2023 7th International Conference on Electronics, Materials Engineering & Nano-Technology (IEMENTech), 2023 Publication	<1 %
71	Submitted to Indian Institute of Technology Student Paper	<1 %
72	Submitted to University of Surrey Student Paper	<1 %
73	Submitted to University of Westminster Student Paper	<1 %

74	Federico Fantechi, Marco Modica. "Learning from the past: a machine-learning approach for predicting the resilience of locked-in regions after a natural shock", Regional Studies, 2022 Publication	<1 %
75	Omkar Rane, Trisha Shishodiya, Roshan Sawant, Anand Godbole. "SignSpeak - Sign Language Interconverter Using CNN-based Approach", 2023 7th International Conference on Electronics, Communication and Aerospace Technology (ICECA), 2023 Publication	<1 %
76	Submitted to The University of Manchester Student Paper	<1 %
77	Submitted to University of Cincinnati Student Paper	<1 %
78	Submitted to University of College Cork Student Paper	<1 %
79	Submitted to University of Northumbria at Newcastle Student Paper	<1 %
80	technodocbox.com Internet Source	<1 %
81	Submitted to Coventry University Student Paper	<1 %

82	Submitted to Higher Education Commission Pakistan Student Paper	<1 %
83	Submitted to Universite Saint Joseph Student Paper	<1 %
84	Submitted to Universiti Malaysia Perlis Student Paper	<1 %
85	Submitted to University of Bedfordshire Student Paper	<1 %
86	Submitted to University of Huddersfield Student Paper	<1 %
87	Submitted to University of Warwick Student Paper	<1 %
88	Submitted to ea Student Paper	<1 %
89	www.codewithc.com Internet Source	<1 %
90	Submitted to Bournemouth University Student Paper	<1 %
91	Submitted to Jose Rizal University Student Paper	<1 %
92	Submitted to Kingston University Student Paper	<1 %
93	Submitted to Middlesex University	

	Student Paper	<1 %
94	Submitted to Nottingham Trent University Student Paper	<1 %
95	Submitted to University of Rwanda Student Paper	<1 %
96	ijircce.com Internet Source	<1 %
97	www.dialog-21.ru Internet Source	<1 %
98	www.irjmets.com Internet Source	<1 %
99	Submitted to Fiji National University Student Paper	<1 %
100	Keren Artiaga, Yang Li, Ercan Engin Kuruoglu, Wai Kin Chan. "Cross-Sign Language Transfer Learning Using Domain Adaptation with Multi-scale Temporal Alignment", Multimedia Tools and Applications, 2023 Publication	<1 %
101	Submitted to Pace University Student Paper	<1 %
102	Submitted to Sim University Student Paper	<1 %

103	Submitted to University of Technology, Sydney Student Paper	<1 %
104	aoyilmaz.medium.com Internet Source	<1 %
105	core.ac.uk Internet Source	<1 %
106	etd.lib.nsysu.edu.tw Internet Source	<1 %
107	open.uct.ac.za Internet Source	<1 %
108	scholarspark.wordpress.com Internet Source	<1 %
109	towardsdatascience.com Internet Source	<1 %
110	www.allresearchjournal.com Internet Source	<1 %
111	www.scaler.com Internet Source	<1 %
112	www.semanticscholar.org Internet Source	<1 %
113	Baydaa Hashim Mohammed, Hasimi Sallehudin, Nurhizam Safie, Mohd Satar, Hamed Dhary Murhg, Shaymaa Abdelghany	<1 %

Mohamed. "Anomaly Detection of Distributed Denial of Service (DDoS) in IoT Network Using Machine Learning", Research Square Platform LLC, 2023
Publication

114	Lecture Notes in Computer Science, 2016. Publication	<1 %
115	Moreira Almeida, Sílvia Grasiella, Frederico Gadelha Guimarães, and Jaime Arturo Ramírez. "Feature extraction in Brazilian Sign Language Recognition based on phonological structure and using RGB-D sensors", Expert Systems with Applications, 2014. Publication	<1 %
116	Submitted to Sheffield Hallam University Student Paper	<1 %
117	Submitted to University of Birmingham Student Paper	<1 %
118	Vijayendra D. Avina, Md Amiruzzaman, Stefanie Amiruzzaman, Linh B. Ngo, M. Ali Akber Dewan. "An AI-Based Framework for Translating American Sign Language to English and Vice Versa", Information, 2023 Publication	<1 %
119	e-space.mmu.ac.uk Internet Source	<1 %

ijarcce.com

120	Internet Source	<1 %
121	iris.poliba.it Internet Source	<1 %
122	link.springer.com Internet Source	<1 %
123	myresearchspace.uws.ac.uk Internet Source	<1 %
124	noexperiencenecessarybook.com Internet Source	<1 %
125	pyimagesearch.com Internet Source	<1 %
126	thesai.org Internet Source	<1 %
127	usermanual.wiki Internet Source	<1 %
128	www.coursehero.com Internet Source	<1 %
129	Ammar Mohamed, Hesham Hefny, mohamed Amin. "A Deep Learning Approach for Gloss Sign Language Translation using Transformer", Journal of Computing and Communication, 2022 Publication	<1 %

130	Regmi, Pratikshya. "Application of UAS Photogrammetry and Geospatial AI Techniques for Palm Tree Detection and Mapping", Texas A&M University - Corpus Christi, 2023 Publication	<1 %
131	Ton Duc Thang University Publication	<1 %

Exclude quotes	On	Exclude matches	Off
Exclude bibliography	On		

3 Acknowledge

Supervised By:

Dr. Nermin Othman

4 Table of Contents

Table of Contents

1 Abstract 2
2 Turnitin Report 3
3 Acknowledge 18
4 Table of Contents 19
5 List of Figures 21
6 List of Tables 24
7 Introduction 25
7.1 Overview 25
7.2 Problem Statement 26
7.3 Scope and Objectives 26
7.4 Report Organization (Structure) 26
7.5 Work Methodology 27
7.6 Work Plan (Gantt chart) 27
8 Related Work (State-of-The-Art) 29
8.1 Background 29
8.2 Literature Survey 29
8.3 Analysis of the Related Work 30
9 Proposed solution 31
9.1 Solution Methodology 31
9.2 Functional/ Non-functional Requirements 32
9.3 Design / Simulation setup 32
10 Implementation 36
10.1 Model Architecture: 36
10.1.1 VGG19 36
10.1.2 SVM 36
10.1.3 Efficient Net 36
10.2 image recognition system: 37
10.2.1 ASL Models: 37
10.2.2 ARSL Model: 53
10.2.3 ASL Realtime: 59
10.2.4 ASL word spelling correction: 65
10.3 Video recognition system: 66

10.3.1 BSL 66
10.4 Video Request: 71
11 Testing and evaluation 73
11.1 Testing 73
11.1.1 Image recognition system 73
11.1.2 Video recognition system 77
11.2 Evaluation 78
11.2.1 Image recognition system 78
11.2.2 Video recognition system 84
12 Results and Discussions 88
12.1 Results 88
12.1.1 Image recognition models 88
12.1.2 Video recognition model 90
12.1.3 Video Request 94
12.2 Discussions 94
12.2.1 image recognition system: 94
12.2.2 Video recognition system: 98
12.2.3 Video Request: 99
13 Conclusions and Future Work 100
13.1 Summary 100
13.2 Future Work 100
14 References 101

5 List of Figures

Figure 1 (pipeline) 31
Figure 2 (Hog feature extraction) 37
Figure 3 (load the images) 38
Figure 4 (X, Y arrays) 38
Figure 5 (split data to train and test) 38
Figure 6 (train machine learning model) 39
Figure 7 (image preprocessing) 39
Figure 8 (train and test split by pytorch) 39
Figure 9 (list of classes) 40
Figure 10 (train and test data loaders) 40
Figure 11 (load pre trained EfficientNetB0) 41
Figure 12 (training structure) 42
Figure 13 (test structure) 43
Figure 14 (training function) 43
Figure 15 (loss function and optimizer) 44
Figure 16 (train Pytorch EfficientNetB0 model) 44
Figure 17 (learning curves) 45
Figure 18 (save the trained model) 46
Figure 19 (list the data) 46
Figure 20 (load data) 47
Figure 21 (X, Y split) 47
Figure 22 (array of X,Y) 47
Figure 23 (train and test split) 48
Figure 24 (label encoder) 48
Figure 25 (load pretrained VGG19) 48
Figure 26 (modified layers for training) 49
Figure 27 (define the optimizer and the hyper parameters) 49
Figure 28 (train VGG19 model) 49
Figure 29 (load data and labels) 50
Figure 30 (train and test split) 50
Figure 31 (train and test data loader) 51
Figure 32 (model creation) 51
Figure 33 (train Pytorch CNN model) 52
Figure 34 (define the data and language) 53
Figure 35 (map creation) 53
Figure 36 (label map) 54
Figure 37 (view number of classes) 54
Figure 38 (print sample with its label) 54
Figure 39 (split the data) 55
Figure 40 (retrieve class label) 56
Figure 41 (load the data) 56
Figure 42 (display prediction) 57
Figure 43 (group train and test data by id) 57
Figure 44 (train test preparation) 58
Figure 45 (model layers) 58

Figure 46 (early stopping function and model training) 59
Figure 47 (data preprocessing) 59
Figure 48 (media pipe hand initialization) 60
Figure 49 (real time prediction) 61
Figure 50 (load trained model) 62
Figure 51 (media pipe hand setup) 62
Figure 52 (images resize) 63
Figure 53 (real time prediction) 64
Figure 54 (word spelling correction) 65
Figure 55 (load data) 66
Figure 56 (label encoder and normalization) 67
Figure 57 (data augmentation) 67
Figure 58 (save data in arrays) 68
Figure 59 (train and test split) 68
Figure 60 (model layers) 69
Figure 61 (model training) 69
Figure 62 (train with base layers) 70
Figure 63 (video retrieval) 71
Figure 64 (SVM testing) 74
Figure 65 (testing function) 75
Figure 66 (load test data) 75
Figure 67 (EfficientNet testing) 76
Figure 68 (predict the labels) 76
Figure 69 (prediction visualization) 77
Figure 70 (confusion matrix) 77
Figure 71 (classification report) 78
Figure 72 (classification report results) 79
Figure 73 (confusion matrix result) 80
Figure 74 (predictions results) 81
Figure 75 (model testing results) 82
Figure 76 (confusion matrix) 83
Figure 77 (Arabic prediction) 84
Figure 78 (confusion matrix results) 85
Figure 79 (confusion matrix results) 86
Figure 80 (classification report) 87
Figure 81 (classification report) 87
Figure 82 (SVM Classification report) 88
Figure 83 (SVM accuracy) 88
Figure 84 (SVM version 2 Classification report) 89
Figure 85 (SVM version 2 accuracy) 89
Figure 86 (training epoch results) 89
Figure 87 (EfficientNet accuracy) 89
Figure 88 (version 2 of EfficientNet epoch) 90
Figure 89 (EfficientNet accuracy) 90
Figure 90 (learning curves) 90
Figure 91 (VGG19 training part 1) 91
Figure 92 (VGG19 training part 2) 91

Figure 93 (VGG19 accuracy) 91
Figure 94 (learning curve results) 92
Figure 95 (VGG19 classification report) 92
Figure 96 (VGG19 confusion matrix) 93
Figure 97 (word selection) 94

6 List of Tables

Table 1 (Related work) 31
Table 2 trials (CNN vs. VGG19) 95
Table 3 final (SVM vs. EfficientNet) 97

7 Introduction

7.1 Overview

Sign language is a form of speech that uses hand shapes, motions for interaction. It is the main method of communication for the deaf and hard-of-hearing. While sign language is rich and flexible, it is little known to those who do not have difficulty hearing. This can create essential communication barriers for the deaf and hard-of-hearing community, limit possibilities for social engagement, work, and learning.

Sign language translation into text has come to be as modern technology with the ability for bringing barriers of communication downward. This technology use machine learning/deep learning algorithms to recognise real-time sign language gestures and translate them to text. This important development allows people who know sign language to communicate with those who don't, increasing more accessibility and equity.

Image recognition and techniques for processing natural languages are usually employed in sign language translation systems. The detection and recognition of sign language gestures from video or image input is the responsibility of image recognition algorithms. These algorithms capture necessary hand features such as position, orientation, and movement and use machine learning models to classify the captured actions.

techniques for processing natural languages are used for converting gestures in sign language into text once they have been recognised. The sequence of recognised gestures is mapped to related phrases or words in the targeted languages. To generate accurate and comfortable translations, techniques for processing natural languages

The advantages of sign language transformation into text go outside just enabling interaction among deaf and hearing people. This technology can also be used to improve accessibility in a variety of settings, such as schools, offices, and public places. Real-time sign language translation can provide live captions for lectures, meetings, and presentations, allowing deaf and hard-of-hearing people to fully participate.

Benefits of sign language translation into text:

- Sign language users benefit from increased communication and social interaction.
- Improved chances for education, employment, and other options for sign language users.
- People who use sign language suffer from less of social isolation and loneliness.

• increasing understanding and awareness of sign language and Deaf Culture.

Challenges of sign language translation into text:

• Sign language is a hard language to know with many different forms and variations.

• Sign language gestures, specifically for real time, can be hard to understand.

• For sign language translation models, there is a lack of high-quality training data.

7.2 Problem Statement

There is a lack of successful interaction among those who utilise sign language and those who do not. This develops obstacles to interaction in many areas of life, including education, employment, and healthcare. The project that I'm working on will involve the creation of a sign language recognition system capable of translating sign language into text in real time. This would enable sign language users to interact more efficiently with others who do not, and vice versa.

The challenges of this solution include:

1. The Sign Language's Complexity.
2. The limited availability of massive amount sign language data datasets.
3. The Sign Language Variance.

7.3 Scope and Objectives

My project's goal is to create a system for the recognition of sign language which can transform sign language into text in real time. The model is going to be capable to recognise a wide range of signs, sign sequences, and entire sentences. The system will also be capable of handling various sign languages.

1. including American Sign Language (ASL).
2. including Arabic Sign Language (ARSL).
3. including Brazilian Sign Language (BSL).

7.4 Report Organization (Structure)

Abstract: The abstract summarizes the report, including the problem statement, proposed solution, and main findings.

Turnitin Report: this part while have the results of turn it in.

Introduction: The Introduction section of the report provides an overview of the problem statement, scope and objectives, report organisation, work methodology, and work plan.

Related Work (State-of-The-Art): The Related Work section provides a re- view of existing research on the problem at hand. This section includes a background on the problem, a literature survey, and an analysis of the related work.

Proposed solution: The Proposed Solution This section includes a solution methodology as well as the solution's functional and non-functional requirements and Design / Simulation setup.

Implementation: this Is the part that has every important part of my code with its description and what it does, and this part consists of model's architectures, image recognition system codes, video recognition system and video request.

Test and evaluation: this part has the testing codes and prediction codes in test part and the results of this testing in evaluation.

Results and Discussions: this is the results of the model and the pipeline like accuracy and learning curves and the discussion part its is the part that has the comparisons and every information related to the models.

Conclusions and Future Work: in the summary part it consists of the Summarise what I have achieved and in future work part it will consist the list of the requirement that I will be working on in the future.

References: The References section provides a list of all sources that were cited in the report

7.5 Work Methodology

- gather a database of sign language videos/images from different places.
- Create a pre-processing pipeline for cleaning and modifying images.
- extract all of features from the videos/images, like as hand shape, direction, and position.
- Create a model to categorise the features and recognise the corresponding sign.
- test the model on a held-out test set to improve accuracy.

7.6 Work Plan (Gantt chart)

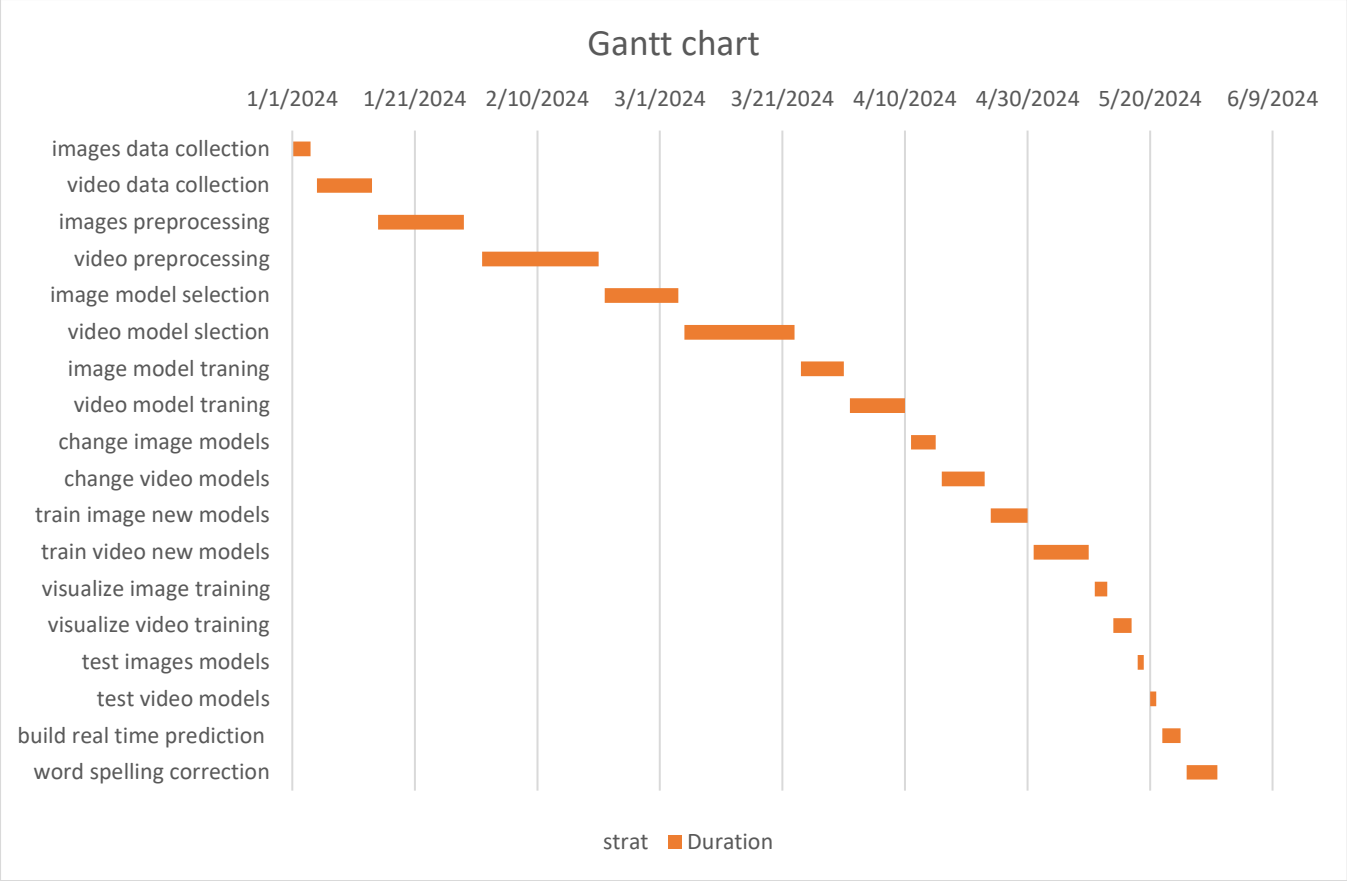
Gantt chart
1/1/2024
1/21/2024
2/10/2024
3/1/2024
3/21/2024
4/10/2024
4/30/2024
5/20/2024
6/9/2024
images data collection
video data collection
images preprocessing
video preprocessing
image model selection
video model slection
image model traning
video model traning
change image models
change video models
train image new models
train video new models
visualize image training
visualize video training
test images models
test video models
build real time prediction
word spelling correction
strat
Duration

8 Related Work (State-of-The-Art)

8.1 Background

Fifth Dimension Technologies (5DT)

is a technologically advanced company that focuses in the development of advanced training solutions through the use of virtual reality (VR) and machine learning models. serving to a wide variety of fields, such as mining, construction, and transportation. Their field of study is in developing complete and realistic training environments that improve operators' security, productivity, and efficacy. 5DT's comprehensive training solutions include simulators, pre-simulators, visualizers, computer-based training (CBT) systems, and walkaround inspection trainers, among other things.

VGG16

It is widely known convolutional neural network (CNN) architecture for recognition of images and classification tasks. VGG16 is made up of 16 layers, 13 of which are convolutional and 3 of which are fully connected. The convolutional layers are used to extract features from the input image, while the layers that are fully connected are used to classify the image. VGG16 is a powerful architecture that has produced state-of-the-art results on a wide range of image recognition tasks.

8.2 Literature Survey

Sign language recognition (SLR) is an evolving field that has the potential to improve communication for deaf and hard-of-hearing people. SLR systems seek to automatically convert sign language gestures into text , allowing for effortless interaction between sign language users and those who do not. Different SLR approaches have been proposed over the last few decades, ranging from traditional machine learning techniques to deep learning models.

In [1] the author proposed this paper use a deep learning model trained on American Sign Language (ASL) is introduced to convert sign language into text in real time. Based on the VGG16 structures, this model use a (CNN) and transfer learning techniques. Through transfer learning techniques, the model's accuracy increased from 94% to 98.7%. The study's main goal is to deal with the obstacles to communication, and the researchers have also created an application that incorporates the deep learning model.

In [2] the author proposed this paper use the segmentation of pictures, feature detection, unsupervised learning, and text to speech technology converting, all this part of the methodology. In unsupervised learned features, the model achieves an accuracy of 78%. The segmentation of captured signed gestures supports in the effective analysis of the images. The K-Nearest Neighbour algorithm is used to classify features, and the greatest match gets transformed to text and speech. Results show that a larger training dataset leads to greater accuracy. They also make use the sensor on the Kinect for gathering data and a MATLAB Speech Synthesiser to generate proper text and speech output.

In [3] the author proposed this paper use data taken from Fifth Dimension Technologies (5DT) gloves and the Australian Sign Language signs dataset to convert gestures into words. The

research compares the efficacy of different machine learning methods, including neural networks, decision tree classifiers, and k-nearest neighbours. The results show that the k-nearest neighbours method performs better at classifying data into words. The research paper highlights the importance of sign language as a way of interaction for deaf and mute people, as well as the need for a translator to help for those who do not understand sign language.

In [4] the author proposed this paper use Deep learning techniques are used to convert sign language to text. Based on hand gestures, the authors created a - 7 - real-time application that predicts and shows the letters in the alphabet. Using the VGG16 model, they were able to achieve an impressive accuracy of 99%. Sign language recognition systems have been developed using methods for image processing and models based on machine learning, such as SVM, Random forestbased approaches, as well as deep learning techniques. The article emphasises the importance of communication for people with hearing problems and the need for effective tools to bridge the communication gap.

In [5] the author proposed this paper use A deep learning method based on a transformer model was used to translate gloss sign language into natural language. It discusses the way conventional sign language translation methods are ineffective and do not have language structure and context. The models proposed in the article exceed other methods in both the translation directions and achieve high scores in evaluation metrics.

In [6] the author proposed this paper use Machine learning methods for sign language recognition including gathering data, pre-processing, categorization, feature extraction, and classification techniques. DeepCon-vLSTM, Graph Convolutional neural Network (GCN), relevance vector machine (RVM), limited condition machine and not clear logic-based, Multiple-layer Perceptron, and Self Organising Maps are also discussed. Among these, the DeepConvLSTM model with da-ta augmentation had a recognition accuracy of 91.1%, while the GCN model had a recognition accuracy of 98.08%, demonstrating the efficacy of these methods.

8.3 Analysis of the Related Work

Papers	Accuracy	Methodology	Disadvantages	Application
[1]	98.70%	Techniques such as CNN and transfer learning	Requires a large training dataset and can be computationally expensive.	Yes
[2]	78%	Image segmentation, feature detection, supervised learning, and technology for text-to-speech	Lower accuracy compared to other methods.	Yes
[3]	95%	Neural networks, decision tree classifiers, and k-nearest neighbors	Limited to recognizing a specific set of gestures and maynot generalize well to other sign languages.	No

[4]	99%	Deep learning techniques	Can only recognize individual letters and not full words or sentences.	Yes
[5]	96.90%	Deep learning approach using a trans-former model	Requires a large amount of training data and may not be suitable for real-time translation.	No
[6]	98.08%	DeepCon-vLSTM, Graph Convolutional neural Net00work (GCN), relevance vector machine (RVM), limited condition machine and not clear logic-based, Multi-layer Perceptron	DeepConvLSTM requires data augmentation, and GCN may be sensitive to noise in the data.	Yes

Table 1 (Related work)

The table compares six different sign language recognition papers in terms of accuracy, methodology, disadvantages, and application. The most accurate paper is [4], which use deep learning techniques. However, the methodology in paper [6] is the most complicated and requires data augmentation. paper [2] has the lowest accuracy and employs image segmentation, feature detection, supervised learning, and text-to-speech technology. This paper, on the other hand, has the simplest methodology and the least computationally expensive. The other four papers use a variety of methodologies, including CNNs, transfer learning, unsupervised learning, and transformer models, and have accuracies ranging from 78% to 98.08%. Some of the disadvantages of these papers involve the require for an enormous training dataset, lack of ability to be applied well to different sign languages, and the necessity for real-time translation.

9 Proposed solution

9.1 Solution Methodology

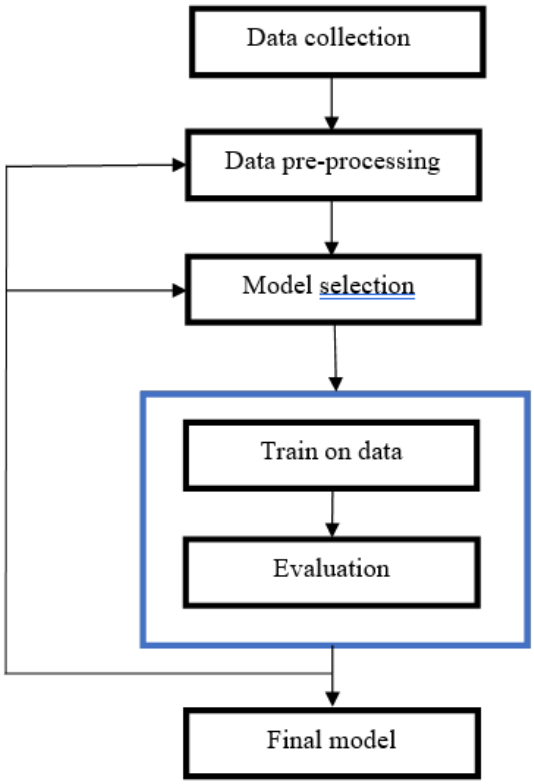

Figure 1 (pipeline)

Data collection:

- gather a dataset of sign language videos/images from different places.

Data pre-processing:

- Create a pre-processing pipeline for cleaning and modifying images.
- extract all of features from the videos/images, like as hand shape, direction, and position.

Model selection:

- Create a machine learning / deep learning model to categorise the features and recognise the corresponding sign.
- test the model on a held-out test set to improve accuracy.

Final model:

- Save Final model.

9.2 Functional/ Non-functional Requirements

Functional Requirements

The system must recognize a wide variety of signs, encompassing both isolated signs and sequences of signs. It should be versatile enough to handle different types of sign languages, particularly American Sign Language (ASL), including both letters and words. Moreover, Real-time operation is essential to ensure smooth and immediate interaction. Additionally, the system should be robust enough to function effectively under various lighting conditions, ensuring consistent performance regardless of environmental changes.

Non-Functional Requirements

Accuracy is paramount for the system to provide reliable translations and interpretations of sign language. Efficiency is also critical, ensuring that the system operates quickly and does not cause delays. The system should be user-friendly, offering an intuitive interface that is easy for users to navigate. Scalability is another important consideration, allowing the system to handle an increasing number of users or signs without compromising performance. Finally, the system should be maintainable, enabling easy updates and improvements to keep up with evolving needs and technological advancements.

9.3 Design / Simulation setup.

1. Overview of the Design Process:

- The design of the sign language recognition system involves several critical stages aimed at creating a robust and efficient solution for translating sign language into text and correcting any spelling errors. The system's primary goal is to provide an accurate operates in real-time. The design process begins with data collection and preprocessing, followed by model training, evaluation, and finally, real-time testing. The ultimate objective is to ensure that the system can recognize and translate sign language gestures accurately and quickly, making it a valuable tool for communication.

2. Requirements and Specifications:

- **Functional Requirements:**

    - The system must recognize and classify a variety of sign language gestures.
    - It should translate recognized gestures into text and correct any spelling errors.
    - The system must operate in real-time, providing immediate feedback.
- **Non-Functional Requirements:**
    - The system must be accurate and efficient.
    - It should be user-friendly and maintainable.

3. Design Methodology:

- The project follows a systematic approach starting from data collection to real-time testing. The methodology includes data preprocessing, model training, and performance evaluation.

4. System Architecture:

The system architecture is designed to integrate all components seamlessly, ensuring robust and efficient performance. The architecture includes the following key components:

- **Data Collection:**
    - Collect images of different sign language gestures from various sources.
    - Label the images accurately to create a comprehensive dataset.
- **Data Preprocessing:**
    - Apply image transformations such as resizing, normalization, and augmentation to enhance the dataset's quality.
    - Ensure that the images are uniform in size and format for consistent model training.
- **Model Training:**
    - Define the model architecture, which includes multiple convolutional layers for feature extraction and fully connected layers for classification.
    - Compile the model with appropriate loss functions and optimizers.
    - Train the model on the pre-processed dataset, adjusting hyperparameters to optimize performance.
- **Evaluation:**
    - Use evaluation metrics such as accuracy, precision, recall, F1-score, and confusion matrix to measure the model’s performance.
    - Plot learning curves to visualize the training and validation process.
- **Real-Time Testing:**
    - Integrate the trained model with OpenCV to capture real-time video frames.
    - Implement real-time predictions, converting captured gestures into text.

    - Use a loop to continuously capture frames and make predictions until the user closes the cv2 window.

- **Word Correction:**
    - Implement a word correction algorithm that uses a dataset of 90K words to correct spelling errors in the predicted text.
    - Integrate the corrected text into the final output, ensuring accurate and meaningful results.

5. Simulation Environment:

- **Software Tools:**
    - Python for programming.
    - TensorFlow/Keras for model training.
    - OpenCV (cv2) for real-time frame capture.
- **Hardware:**
    - Standard PC or laptop with webcam for real-time testing.
- **Libraries:**
    - NumPy, Pandas, Matplotlib for data manipulation and visualization.
    - Scikit-learn for evaluation metrics.

6. Setup Procedures:

- **Data Collection:**
    - Gather and label sign language images.
- **Preprocessing:**
    - Apply image transformations like resizing, normalization, and augmentation.
- **Model Training:**
    - Define the model architecture, compile the model, and train on the pre-processed data.
- **Evaluation:**
    - Plot learning curves, generate confusion matrices, and classification reports.
- **Real-Time Testing Setup:**
    - Install and configure OpenCV for capturing video frames.
    - Integrate the trained model with cv2 to make real-time predictions.
- **Word Correction Integration:**

    - Implement the word correction algorithm and integrate it with the real-time prediction system.

7. Validation and Testing:

Validation and testing are crucial to ensure the system performs reliably under real-world conditions:

- **Validation:**
    - Use a separate validation dataset to fine-tune the model and avoid overfitting.
    - Validate the model's performance using various evaluation metrics.
- **Testing:**
    - Conduct extensive testing in real-time scenarios to ensure the system's robustness and accuracy.
    - Test the system under different lighting conditions and backgrounds to evaluate its adaptability.
- **Evaluation Metrics:**
    - Measure the system's performance using metrics such as accuracy, precision, recall, F1-score, and confusion matrix.

8. Expected Outcomes:

- **Model Performance:**
    - High accuracy in recognizing sign language gestures.
    - Effective real-time performance with minimal latency.
- **User Experience:**
    - Seamless interaction with the system providing immediate and accurate text output.
    - Corrected text output with minimal spelling errors.

9. Challenges and Mitigations:

Identifying and addressing potential challenges is crucial to ensure the system's success:

- **Challenges:**
    - Variability in lighting conditions and backgrounds during real-time testing.
    - Handling different user sign styles and speeds, which may affect recognition accuracy.
- **Mitigations:**
    - Implement robust preprocessing techniques to normalize input data and reduce variability.

- Use data augmentation to train the model on a diverse set of images, improving its generalization capability.
- Continuously update and refine the word correction algorithm to enhance accuracy and reliability.

This comprehensive approach ensures that the design and simulation setup for the sign language recognition system is thorough, well-structured, and capable of achieving the desired outcomes, providing a reliable and efficient tool for communication.

10 Implementation

10.1 Model Architecture:

10.1.1 VGG19

VGG19 is a deep convolutional neural network (CNN) model known for its simplicity and effectiveness in image classification tasks. The architecture consists of 19 layers, including 16 convolutional layers grouped into five blocks. Each block comprises multiple convolutional layers with small 3x3 filters, followed by a max-pooling layer that reduces the spatial dimensions. The number of filters doubles after each block, starting from 64 and reaching up to 512 in the deeper layers. After the convolutional blocks, VGG19 includes three fully connected (FC) layers: two with 4096 units each and a final layer with 1000 units for classification. Each convolutional and fully connected layer is followed by a ReLU activation function, with a softmax layer at the end for outputting class probabilities. This architecture, despite its depth, maintains simplicity and consistency in its design.

10.1.2 SVM

Support Vector Machine (SVM) is a powerful supervised machine learning algorithm used for both classification and regression tasks. The primary goal of SVM is to find the optimal hyperplane that maximally separates different classes in the feature space. SVM works by mapping input data into a high-dimensional feature space where a linear separator can be applied. This is achieved using kernel functions such as linear, polynomial, radial basis function (RBF), and sigmoid, which enable SVM to handle non-linear classification problems effectively. SVM is known for its robustness and effectiveness in high-dimensional spaces, making it a popular choice for tasks such as text classification, image recognition, and bioinformatics.

10.1.3 Efficient Net

Efficient Net is a state-of-the-art convolutional neural network (CNN) architecture designed for image recognition and classification tasks. Developed by Google, Efficient Net scales the model dimensions—depth, width, and resolution—using a compound scaling method, which ensures a more efficient use of resources while maintaining high performance. The base model, EfficientNet-B0, serves as the foundation, and subsequent models (B1 to B7) build upon it by increasing complexity and computational requirements. Efficient Net has demonstrated superior accuracy and efficiency on a variety of benchmarks, outperforming previous models like ResNet and Inception in

terms of both accuracy and computational cost, making it a popular choice for deploying in real-world applications where computational resources are limited.

My implementation consists of 3 main tasks each task has its own sup tasks each sup tasks has its own trials first we going to start with images part.

We used in this part 2 datasets first dataset is for Arabic sign language (ARSL) it consist of 32 class and 50000 images and second dataset consist of 4 merged dataset for American sign language (ASL) it now consist of 29 class and 367528 images.

First sub part from images part is image recognition system and the second part is real time prediction, and third part is word correction.

10.2 image recognition system:

10.2.1 ASL Models:

consist of CNN, Efficient Net, SVM, VGG19 the (Efficient Net and SVM) are the final models.

10.2.1.1 SVM.

Pre-processing:

```
# Function to preprocess the image and extract HOG features
def preprocess_image(image_path):
    image = io.imread(image_path) # Read the image
    image = transform.resize(image, (64, 64)) # Resize to 64x64
    image = color.rgb2gray(image) # Convert to grayscale
    hog_features = feature.hog(image, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
    return hog_features
```

Figure 2 (Hog feature extraction)

function preprocess image that preprocesses an image and extracts Histogram of Oriented Gradients (HOG) features for further analysis or model training. The function takes the path to an image as input and performs several operations: it reads the image from the specified path using io.imread, resizes the image to 64x64 pixels using transform.resize, converts the image to grayscale with color.rgb2gray, and finally, extracts HOG features using feature.hog with specific parameters for pixels per cell and cells per block. The HOG features are then returned by the function, providing a processed feature set suitable for machine learning model.

```
# Iterate through each folder in the training data directory
for folder in os.listdir(train_data_dir):
    folder_path = os.path.join(train_data_dir, folder)
    if os.path.isdir(folder_path):
        label = folder
        for image_name in os.listdir(folder_path):
            image_path = os.path.join(folder_path, image_name)
            # Preprocess the image and extract HOG features
            features = preprocess_image(image_path)
            hog_features.append(features)
            labels.append(label)
```

Figure 3 (load the images)

code iterates through each folder in the specified training data directory, where each folder represents a different label for the training images. For each folder, the code constructs the folder path and checks if it is a directory. If it is, it assigns the folder name as the label for the images contained within. The code then iterates through each image file in the folder, constructs the full path for each image, and preprocesses the image by extracting its HOG features using the previously defined preprocess image function. The extracted HOG features are appended to the hog features list, and the corresponding labels are appended to the labels list. This process systematically prepares the feature set and labels required for training a machine learning model on the given dataset.

```
# Convert lists to numpy arrays
X = np.array(hog_features)
y = np.array(labels)
```

Figure 4 (X, Y arrays)

code snippet converts lists of features and labels into NumPy arrays using the np.array function. Specifically, it transforms the hog_features list, which contains the extracted Histogram of Oriented Gradients (HOG) features for each image, into a NumPy array X. Similarly, it converts the labels list, which holds the corresponding labels for each image, into a NumPy array y. Converting these lists into NumPy arrays is essential for efficient numerical computations and compatibility with machine learning libraries, which typically require input data to be in array format for training and evaluation processes.

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Figure 5 (split data to train and test)

code uses the train_test_split function from the scikit-learn library to divide the dataset into training and testing sets. The function takes the features (X) and labels (y) as input and splits them into four subsets: X_train, X_test, y_train, and y_test. The test_size parameter is set to 0.2, meaning 20% of the data is allocated for testing, and the remaining 80% is used for training the model. The random_state parameter is set to 42 to ensure the split is reproducible and consistent across

different runs. This step is crucial for evaluating the model's performance on unseen data and preventing overfitting.

Model training:

```
# Train the model on the new data
model.fit(X_train, y_train)
```

Figure 6 (train machine learning model)

code shows the step where the model is trained on the training data. Using the fit method of the model object, the model is trained with the features (X_train) and corresponding labels (y_train).

10.2.1.2 Efficient Net.

Preprocessing:

```
import torchvision
from torchvision import datasets
import torchvision.transforms as transforms

print(torchvision.__version__)

# Transform image
data_transform = transforms.Compose([
    transforms.Resize(size = (128, 128)),        # Resize our images to 128X128
    transforms.ToTensor()
])


transformed_data = data_transform(img)
transformed_data, transformed_data.shape, transformed_data.dtype
```

Figure 7 (image preprocessing)

The is creating the data_transform pipeline, created using transforms.Compose. This pipeline includes two transformations: resizing the image to 128x128 pixels with transforms.Resize and converting the image to a tensor with transforms.ToTensor. The transformed data is then applied to an image, and the code outputs the transformed data, its shape, and data type. This process is essential for preparing images in a consistent format required for input into deep learning models.

```
train_size = int(0.8 * len(source_data))
test_size = len(source_data) - train_size

# Split the dataset into train and test sets
train_data, test_data = random_split(source_data, [train_size, test_size])

print(f"Train set size: {len(train_data)}")
print(f"Test set size: {len(test_data)}")
```

Figure 8 (train and test split by pytorch)

This code split a dataset into training and testing sets using PyTorch. First, it calculates the sizes for the training and testing sets, setting the training set size to 80% of the total dataset (train_size) and the testing set size to the remaining 20% (test_size). The dataset (source_data) is then split into training and testing subsets using the random_split function, which randomly allocates data points to each subset based on the calculated sizes. Finally, it prints the sizes of the training and testing sets to confirm the split. This process ensures that the model has separate data for training and evaluation, helping to assess its performance on unseen data.

```
# Get class names as list and dict
class_names = source_data.classes
class_dict = source_data.class_to_idx
class_dict
```

Figure 9 (list of classes)

The code extracts the class names and their corresponding indices from the dataset (source_data). The class_names variable stores the list of class names, which are the labels used to categorize the data. The class_dict variable stores a dictionary that maps each class name to its respective index in the dataset, accessible via source_data.class_to_idx. This mapping is essential for understanding the relationship between the class labels and their numeric representations used by the model during training and prediction.

```
# Train and test dataloaders
from torch.utils.data import DataLoader

BATCH_SIZE = 192
train_dataloader = DataLoader(dataset = train_data,
                              batch_size = BATCH_SIZE,
                              num_workers = 1,
                              shuffle = True)

test_dataloader = DataLoader(dataset = test_data,
                             batch_size = BATCH_SIZE,
                             num_workers = 1,
                             shuffle = False)

len(train_dataloader), len(test_dataloader)
```

Figure 10 (train and test data loaders)

The code set up DataLoaders for training and testing datasets in PyTorch, facilitating efficient data loading and processing. The batch size is set to 192, and the DataLoaders are created using the DataLoader class from torch.utils.data. For the train_dataloader, shuffling is enabled (shuffle=True) to ensure the data is randomly mixed during training, which helps in reducing overfitting. For the test_dataloader, shuffling is disabled (shuffle=False) to maintain the original order of the test data. The number of worker threads for data loading is set to 1 (num_workers=1) for both DataLoaders. The lengths of the train_dataloader and test_dataloader are then printed to confirm the setup. This setup is crucial for feeding data into the model in manageable batches, ensuring efficient training and evaluation processes.

<u>Setup the Model and training functions:</u>

```python
from efficientnet_pytorch import EfficientNet
from torchinfo import summary

num_classes = len(class_names)    # Number of classes

# Instantiate the EfficientNet model
model = EfficientNet.from_pretrained('efficientnet-b0').to(device)
model._fc = nn.Linear(model._fc.in_features, num_classes)

# Wrap the real model with DataParallel to use multiple GPUs
if torch.cuda.device_count() > 1:
    print(f"Let's use {torch.cuda.device_count()} GPUs!")
    model = nn.DataParallel(model, device_ids=[0, 1])

# Move the real model to GPU(s)
model.to(device)

summary(model=model,
        input_size=(192, 3, 128, 128),
        col_names=["input_size", "output_size", "num_params", "trainable"],
        col_width=20,
        row_settings=["var_names"])
```

Figure 11 (load pre trained EfficientNetB0)

The code set up and configures an EfficientNet model for training. First determining the number of classes from class_names. The EfficientNet model (efficientnet-b0) is instantiated using a pre-trained version and moved to the specified device (GPU). The final fully connected layer (_fc) of the model is modified to match the number of output classes. the model is wrapped with DataParallel to enable parallel processing across GPUs, enhancing computational efficiency. The model is then moved to the designated device. The summary function from torchinfo provides a detailed overview of the model architecture, including input size, output size, number of parameters, and which parameters are trainable. This setup is crucial for preparing the model for efficient and effective training on a possibly multi-GPU environment.

```python
def train_step(model, dataloader, loss_fn, optimizer):
    model.train()
    train_loss, train_acc = 0, 0

    # Loop through data loader data batches
    for batch, (X, y) in enumerate(dataloader):
        X, y = X.to(device), y.to(device)
        y_pred = model(X)

        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Calculate and accumulate accuracy metric across all batches
        y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim=1)
        train_acc += (y_pred_class == y).sum().item()/len(y_pred)

    # Adjust metrics to get average loss and accuracy per batch
    train_loss = train_loss / len(dataloader)
    train_acc = train_acc / len(dataloader)
    return train_loss, train_acc
```

Figure 12 (training structure)

train_step function that performs one epoch of training for a given model using a data loader, loss function, and optimizer. The function starts by setting the model to training mode with model.train() and initializes the cumulative training loss and accuracy to zero. It then iterates over batches of data from the data loader. For each batch, the input features (X) and labels (y) are moved to the specified device (GPU). The model makes predictions (y_pred), and the loss is calculated using the specified loss function. The optimizer's gradients are reset with optimizer.zero_grad(), and the loss is backpropagated. The optimizer updates the model parameters with optimizer.step(). Accuracy is calculated by comparing the predicted class (obtained using torch.argmax) with the true labels, and both loss and accuracy are accumulated over all batches. Finally, the average loss and accuracy per batch are computed and returned. This function encapsulates the essential steps for training a neural network, ensuring that the model learns from the data efficiently.

```python
def test_step(model, dataloader, loss_fn):
    model.eval()
    test_loss, test_acc = 0, 0

    # Turn on inference
    with torch.inference_mode():
        for batch, (X, y) in enumerate(dataloader):
            X, y = X.to(device), y.to(device)

            test_pred_logits = model(X)
            loss = loss_fn(test_pred_logits, y)
            test_loss += loss.item()

            # Calculate and accumulate accuracy
            test_pred_labels = test_pred_logits.argmax(dim=1)
            test_acc += ((test_pred_labels == y).sum().item()/len(test_pred_labels))

    # Adjust metrics to get average loss and accuracy per batch
    test_loss = test_loss / len(dataloader)
    test_acc = test_acc / len(dataloader)
    return test_loss, test_acc
```

Figure 13 (test structure)

test_step function that evaluates a given model's performance on a test dataset using a data loader and a loss function. The function starts by setting the model to evaluation mode with model.eval() and initializing cumulative test loss and accuracy to zero. It uses torch.inference_mode() to disable gradient computation, optimizing inference speed and reducing memory usage. The function then iterates over batches of data from the data loader. For each batch, the input features (X) and labels (y) are moved to the specified device (GPU). The model makes predictions (test_pred_logits), and the loss is calculated using the specified loss function. The cumulative test loss is updated accordingly. Accuracy is calculated by comparing the predicted class labels (obtained using argmax) with the true labels, and this accuracy is accumulated over all batches. Finally, the average loss and accuracy per batch are computed and returned. This function encapsulates the essential steps for evaluating a neural network's performance ensuring accurate and efficient assessment of the model on unseen data.

```python
def train(model, train_dataloader, test_dataloader, optimizer, loss_fn, epochs):

  # Create a empty results dictionary
  results = {"train_loss": [],
             "train_acc": [],
             "test_loss":[],
             "test_acc": []}

  # Loop through training and testing steps for no of epochs
  for epoch in tqdm(range(epochs)):
    train_loss, train_acc = train_step(model = model,
                                       dataloader = train_dataloader,
                                       optimizer = optimizer,
                                       loss_fn = loss_fn)

    test_loss, test_acc = test_step(model = model,
                                    dataloader = test_dataloader,
                                    loss_fn = loss_fn)

    print(f'Epoch: {epoch} | Train loss: {train_loss:.4f} - Train acc: {(train_acc*100):.2f}% -  Test loss: {test_loss:.4f} - Test acc: {(test_acc*100):.2f}%')

    # Update results dictionary
    results["train_loss"].append(train_loss)
    results["train_acc"].append(train_acc)
    results["test_loss"].append(test_loss)
    results["test_acc"].append(test_acc)

  return results
```

Figure 14 (training function)

train function that manages the training and testing of a model over a specified number of epochs. The function takes as inputs the model, training, and testing dataloaders, optimizer, loss function, and the number of epochs. It initializes a dictionary result to store the training and testing loss and

accuracy for each epoch. Within a loop that iterates over the number of epochs, the train_step function is called to perform a single training epoch, updating the model weights and calculating training loss and accuracy. Similarly, the test_step function is called to evaluate the model on the test dataset, computing the test loss and accuracy. The results of each epoch are printed and appended to the results dictionary. This function effectively manages the training and evaluation process, providing a comprehensive record of model performance over time.

Model training:

```
# Setup loss functions and optimizer
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(params = model.parameters(), lr=0.001)
```

Figure 15 (loss function and optimizer)

The code set up the loss function and optimizer for training a neural network model. The loss function is defined as CrossEntropyLoss, which is commonly used for multi-class classification tasks and measures the difference between the predicted and true class distributions. The optimizer is defined as the Adam optimizer torch.optim.Adam, which is an adaptive learning rate optimization algorithm designed to handle sparse gradients on noisy problems. It is initialized with the model's parameters and a learning rate of 0.001. This setup is crucial for guiding the model's learning process during training, as the loss function provides the feedback on prediction accuracy, and the optimizer adjusts the model parameters to minimize the loss.

```
EPOCHS = 5

# Start the timer
from timeit import default_timer as timer
start_time = timer()

# Train model
model_results = train(model = model,
                      train_dataloader = train_dataloader,
                      test_dataloader = test_dataloader,
                      optimizer = optimizer,
                      loss_fn = loss_fn,
                      epochs = EPOCHS)

# End the timer and print duration of training
end_time = timer()
print(f'Total Train Time: {end_time - start_time:.3f} seconds')
```

Figure 16 (train Pytorch EfficientNetB0 model)

The provided code snippet sets up the training process for a model over a specified number of epochs and measures the total training time using the default_timer from the timeit module. It defines EPOCHS as 5, indicating the number of training cycles. The timer is started by calling timer(), and the train function is executed, which handles the training and testing of the model using the provided dataloaders, optimizer, loss function, and the number of epochs. After training completes, the timer is stopped, and the total duration of the training process is calculated and printed. This snippet effectively integrates timing into the training workflow, providing insights into the computational efficiency and duration of the training process.

Visualization:

```python
def plot_loss_curves(results):
    """ Plots training curves of a results dictionary """

    # Get results dictionary (training and test)
    loss = results['train_loss']
    test_loss = results['test_loss']
    accuracy = results['train_acc']
    test_accuracy = results['test_acc']

    epochs = range(len(results['train_loss']))

    plt.figure(figsize=(15, 4))

    # Plot loss
    plt.subplot(1, 2, 1)
    plt.plot(epochs, loss, label='train_loss')
    plt.plot(epochs, test_loss, label='test_loss')
    plt.title('Loss')
    plt.xlabel('Epochs')
    plt.legend()

    # Plot accuracy
    plt.subplot(1, 2, 2)
    plt.plot(epochs, accuracy, label='train_accuracy')
    plt.plot(epochs, test_accuracy, label='test_accuracy')
    plt.title('Accuracy')
    plt.xlabel('Epochs')
    plt.legend();

plot_loss_curves(model_results)
```

Figure 17 (learning curves)

The code defines a function plot_loss_curves that generates and displays training and validation loss and accuracy curves from a results dictionary. The dictionary contains training and testing loss (train_loss, test_loss) and accuracy (train_acc, test_acc) recorded over each epoch. The function begins by extracting these values and setting up the range of epochs. It then creates a figure with a specified size and plots two subplots: one for loss and one for accuracy. The loss subplot shows both training and testing loss over epochs, while the accuracy subplot displays training and testing accuracy. Each subplot is titled and labelled appropriately, and legends are added for clarity. The plot_loss_curves function is then called with model_results to visualize the model's performance over the training period. This visualization helps in analysing how well the model is learning and generalizing to new data.

Saving the model:

```python
MODEL_PATH = Path("models")
MODEL_PATH.mkdir(parents=True,
                 exist_ok=True)

# Create model save path
MODEL_NAME = "ASL_HandSignLang_EfficientNetB0_2.pth"
MODEL_SAVE_PATH = MODEL_PATH / MODEL_NAME

# Save the model state dict
print(f"Saving model to: {MODEL_SAVE_PATH}")
torch.save(obj=model.state_dict(),
           f=MODEL_SAVE_PATH)
```

Figure 18 (save the trained model)

The code is a directory for saving a trained model and then saves the model's state dictionary to a specified file path. It begins by defining MODEL_PATH as a Path object pointing to a directory named "models" and creates this directory if it doesn't exist (mkdir with parents=True and exist_ok=True). The model's save path is constructed by combining MODEL_PATH with MODEL_NAME, which is defined as "ASL_HandSignLang_EfficientNetB0_2.pth". The model's state dictionary, which contains all the parameters of the model, is saved to this path using torch.save. A message is printed to indicate the save location. This process ensures that the trained model can be saved and later loaded for inference or further training, preserving the model's learned parameters.

10.2.1.3 VGG19

<u>Preprocessing</u>

```python
# Change the directory path
dataset_path = r"C:\Users\Administrator\Desktop\GP\datasets\American"
# List all directories in the dataset path, assuming each directory represents a class
classes = [d for d in os.listdir(dataset_path) if os.path.isdir(os.path.join(dataset_path, d))]
classes
```

Figure 19 (list the data)

The code sets the path to a dataset directory and lists all subdirectories within it, with each subdirectory representing a different class. The dataset_path variable is assigned the path to the dataset directory ("C:\Users\Administrator\Desktop\GP\datasets\American"). The classes list is then populated by iterating through the contents of the specified dataset directory using os.listdir. It includes only those entries that are directories, identified by checking with os.path.isdir. This results in a list of class names corresponding to the subdirectories within the dataset path, which can be used for labelling and organizing the data for further processing or model training.

```
data = []
# Iterate over each class directory
for c in classes:
    class_path = os.path.join(dataset_path, c)  # Full path to the class directory
    # Process each image in the class directory
    for img in os.listdir(class_path):
        img_path = os.path.join(class_path, img)  # Full path to the image
        im = cv2.imread(img_path)  # Read the image
        if im is not None:  # Check if the image was correctly loaded
            im = cv2.resize(im, (224, 224))  # Resize the image
            data.append([im, c])  # Append the image and its class label to the data list
```

Figure 20 (load data)

The code processes images from a dataset where each subdirectory represents a different class. It initializes an empty list data to store the processed images and their labels. For each class in the classes list, it constructs the full path to the class directory. It then iterates over each image file in the class directory, constructs the full path to the image, and reads the image using cv2.imread. If the image is successfully loaded, it is resized to 224x224 pixels using cv2.resize. The resized image and its corresponding class label are appended as a tuple to the data list. This process systematically reads, processes, and stores images along with their labels, preparing the data for subsequent steps such as training a machine learning model.

```
X=[]
Y=[]
for x,y in data:
    X.append(x)
    Y.append(y)
```

Figure 21 (X, Y split)

The code separates the processed image data and their corresponding labels into two separate lists, `X` and `Y`. It initializes these lists as empty. Then, it iterates over the `data` list, which contains tuples of images and their class labels. For each tuple `(x, y)` in `data`, the image `x` is appended to the `X` list, and the label `y` is appended to the `Y` list. This operation organizes the dataset into two distinct lists: one for the feature data (images) and one for the target labels, facilitating subsequent steps such as data transformation, model training, or evaluation.

```
x=np.array(X)
y=np.array(Y)
print(x.shape,y.shape)
```

Figure 22 (array of X,Y)

The provided code snippet converts the lists of images (X) and their corresponding labels (Y) into NumPy arrays for efficient numerical computation and compatibility with libraries. The np.array function is used to convert X into the array x and Y into the array y. After the conversion, the shapes of the resulting arrays x and y are printed using print (x.shape, y.shape). This step ensures that the data is in the correct format and dimensions, facilitating further processing steps such as data normalization, training, and evaluation in machine learning workflows.

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=4, stratify=y)
```

Figure 23 (train and test split)

The code uses the train_test_split function from the scikit-learn library to split the dataset into training and testing sets. The arrays x (features) and y (labels) are divided such that 30% of the data is allocated to the test set (x_test, y_test), and the remaining 70% is allocated to the training set (x_train, y_train). The random_state parameter is set to 4 to ensure reproducibility of the split, and the stratify=y parameter ensures that the class distribution in the training and test sets matches the original dataset. This stratified split is crucial for maintaining balanced class representation, which helps improve the reliability of the model evaluation.

```
# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Fit label encoder and transform labels
y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)
```

Figure 24 (label encoder)

The code demonstrates the use of the LabelEncoder from the scikit-learn library to encode categorical labels into numeric form, which is required for models. The LabelEncoder is initialized and then fitted to the training labels (y_train), transforming them into numerical format and storing the result in y_train_encoded. The same encoder is then used to transform the test labels (y_test), ensuring consistency in the label encoding. The transformed test labels are stored in y_test_encoded. This process converts the categorical labels into a format that can be utilized by machine learning models, facilitating accurate training and evaluation.

<u>Build VGG19 model:</u>

```
model_vgg19=tf.keras.applications.vgg19.VGG19(
    input_shape=None,
    include_top=True,
    weights='imagenet',
    input_tensor=None,
    pooling=None,
    classes=1000,
    classifier_activation='softmax'
    #callbacks=[checkpointer]
)
model_vgg19.summary()
```

Figure 25 (load pretrained VGG19)

The code initializes a pre-trained VGG19 model using TensorFlow's Keras API from the tf.keras.applications module. The model is set up with several parameters: input_shape is unspecified (allowing for variable input sizes), include_top=True includes the fully connected layers at the top of the network, weights='imagenet' loads pre-trained weights from the ImageNet dataset, input_tensor is unspecified, pooling is unspecified, classes=1000 indicates the number of output classes, and classifier_activation='softmax' specifies the activation function for the classifier. The line model_vgg19.summary() prints a summary of the model's architecture, providing details about each

layer, its output shape, and the number of parameters. This setup allows for leveraging a robust, pre-trained VGG19 model for various image classification tasks.

```
base_input4=model_vgg19.input
base_output4=model_vgg19.layers[-2].output

final_output4 = layers.Dense(240)(base_output4)
final_output4 = layers.Activation('relu')(final_output4)
final_output4 = layers.Dense(128)(final_output4)
final_output4 = layers.Activation('relu')(final_output4)
final_output4 = layers.Dense(64)(final_output4)
final_output4 = layers.Activation('relu')(final_output4)
final_output4 = layers.Dense(37, activation='softmax')(final_output4)
```

Figure 26 (modified layers for training)

The code customizes the VGG19 model by adding new fully connected (dense) layers to the existing architecture. It begins by referencing the model's input (`base_input4`) and the output of the penultimate layer (`base_output4`). Then, several new dense layers are sequentially added to `base_output4`, each followed by a ReLU activation function. Specifically, the layers are as follows: a dense layer with 240 units, a dense layer with 128 units, and a dense layer with 64 units, each followed by a ReLU activation. Finally, the last dense layer has 37 units with a softmax activation function, indicating that the model is designed for a classification task with 37 classes. This setup effectively creates a new output head for the VGG19 model, tailored to the specific number of classes in the new dataset while leveraging the pre-trained features from the VGG19 backbone.

```
new_model_vgg19 = tf.keras.Model(inputs = base_input4, outputs = final_output4)


new_model_vgg19.compile(loss="sparse_categorical_crossentropy", optimizer = "SGD", metrics = ["accuracy"])
```

Figure 27 (define the optimizer and the hyper parameters)

The code defines and compiles a new model based on the customized VGG19 architecture. It begins by creating a new model instance new_model_vgg19 using tf.keras.Model, specifying the base_input4 as the input layer and final_output4 as the output layer, effectively combining the pre-trained VGG19 model with the newly added dense layers. The new model is then compiled with the compile method, setting the loss function to sparse_categorical_crossentropy suitable for classification tasks with integer labels. The optimizer is specified as Stochastic Gradient Descent (SGD), and the evaluation metric is set to accuracy. This configuration prepares the model for training, ensuring it uses the specified loss function, optimization algorithm, and performance metric to learn from the data effectively.

<u>model training:</u>

```
# Using a smaller subset of data
subset_size = 1000  # adjust this number based on your dataset size and available memory
result_V = new_model_vgg19.fit(x_train[:subset_size], y_train_encoded[:subset_size], epochs=10, batch_size=64, validation_data=(x_test[:subset_size], y_test_encoded[:subset_size]))
```

Figure 28 (train VGG19 model)

The code trains the newly defined VGG19 model on a smaller subset of the dataset to accommodate memory constraints or speed up the training process. It sets the subset_size to 1000, meaning only

the first 1000 samples from the training and testing datasets will be used. The fit method is called on new_model_vgg19, training the model on x_train and y_train_encoded limited to the subset size. The training runs for 10 epochs with a batch size of 64. The validation data is similarly limited to the first 1000 samples of x_test and y_test_encoded. This approach allows for faster iteration and testing of the model on a smaller, manageable portion of the data while still providing insights into its performance.

10.2.1.4 CNN:

Preprocessing:

```
folder_names = os.listdir("/kaggle/input/aslamerican-sign-language-aplhabet-dataset/ASL_Alphabet_Dataset/asl_alphabet_train")
labels = []
data_images = []

for index, folder in enumerate(folder_names):
    image_files = os.listdir(f"/kaggle/input/aslamerican-sign-language-aplhabet-dataset/ASL_Alphabet_Dataset/asl_alphabet_train/{folder}/")
    for image_file in image_files:
        image_path = f"/kaggle/input/aslamerican-sign-language-aplhabet-dataset/ASL_Alphabet_Dataset/asl_alphabet_train/{folder}/{image_file}"
        image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
        image = cv2.resize(image, (32, 32))
        image = image / 255.0
        data_images.append([image])
        labels.append(index)

data_images = np.array(data_images, dtype=np.float32)
labels = np.array(labels, dtype=np.float32)
```

Figure 29 (load data and labels)

The code processes a dataset of American Sign Language (ASL) alphabet images from a specified directory and prepares it for training. It first lists all folder names in the dataset directory, where each folder represents a different sign. It initializes empty lists data_images and labels to store the processed images and their corresponding labels. The code then iterates through each folder and its images. For each image, it constructs the image path, reads the image in grayscale using cv2.imread, resizes it to 32x32 pixels, and normalizes the pixel values by dividing by 255. The processed image is appended to data_images, and the corresponding label (folder index) is appended to labels. Finally, data_images and labels are converted to NumPy arrays with a data type of float32. This prepares the data in a format suitable for input into a machine learning model for further processing or training.

```
features_train, features_test, labels_train, labels_test = train_test_split(data_images, labels, test_size=0.2, random_state=42)
featuresTrain = torch.from_numpy(features_train)
labelsTrain = torch.from_numpy(labels_train).type(torch.LongTensor)
featuresTest = torch.from_numpy(features_test)
labelsTest = torch.from_numpy(labels_test).type(torch.LongTensor)
```

Figure 30 (train and test split)

The code splits the processed ASL alphabet dataset into training and testing sets and converts them into PyTorch tensors. Initially, the train_test_split function from scikit-learn is used to split data_images and labels into training (features_train, labels_train) and testing sets (features_test, labels_test), with 20% of the data allocated for testing (test_size=0.2) and a fixed random state for reproducibility (random_state=42). Each subset is then converted from NumPy arrays to PyTorch tensors using torch.from_numpy, ensuring compatibility with PyTorch's data processing pipeline. The labels are cast to LongTensor

type, which is required for certain PyTorch operations, particularly those involving categorical data. This setup prepares the data for training and evaluating a PyTorch model.

```
train = DataLoader(torch.utils.data.TensorDataset(featuresTrain, labelsTrain), batch_size=100, shuffle=False)
test = DataLoader(torch.utils.data.TensorDataset(featuresTest, labelsTest), batch_size=100, shuffle=False)
```

Figure 31 (train and test data loader)

The code sets up DataLoaders for the training and testing datasets. The training DataLoader (train) and the testing DataLoader (test) are created by passing TensorDataset instances, which combine the features and labels tensors (featuresTrain, labelsTrain for training and featuresTest, labelsTest for testing). Both DataLoaders are configured with a batch size of 100, meaning the data will be processed in batches of 100 samples at a time. The shuffle=False parameter indicates that the data will not be shuffled during loading, maintaining the original order of samples. These DataLoaders facilitate efficient data loading and iteration during the model training and evaluation phases.

<u>Model creation:</u>

```
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()
        self.cnn1 = nn.Conv2d(1, 16, 5, 1, 0)
        self.relu1 = nn.ReLU()
        self.maxpool1 = nn.MaxPool2d(2)
        self.cnn2 = nn.Conv2d(16, 32, 5, 1, 0)
        self.relu2 = nn.ReLU()
        self.maxpool2 = nn.MaxPool2d(2)
        self.fc1 = nn.Linear(32 * 5 * 5, 29)

    def forward(self, x):
        x = self.maxpool1(self.relu1(self.cnn1(x)))
        x = self.maxpool2(self.relu2(self.cnn2(x)))
        x = x.view(x.size(0), -1)
        x = self.fc1(x)
        return x

model = CNNModel()
error = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

Figure 32 (model creation)

The code defines a convolutional neural network (CNN) model class, CNNModel. The model's architecture is specified in the __init__ method, which includes two convolutional layers (self.cnn1 and self.cnn2), each followed by a ReLU activation (self.relu1 and self.relu2) and a max-pooling layer (self.maxpool1 and self.maxpool2). The first convolutional layer has

16 filters with a kernel size of 5, and the second has 32 filters with a kernel size of 5. The fully connected layer (self.fc1) maps the flattened output of the final max-pooling layer to 29 output classes, corresponding to the ASL alphabet. The forward method defines the forward pass of the model, sequentially applying the convolutional, activation, and pooling layers, flattening the output, and passing it through the fully connected layer. The model is instantiated, and a cross-entropy loss function (error) and stochastic gradient descent (SGD) optimizer (optimizer) with a learning rate of 0.1 are defined. This setup prepares the CNN model for training on the ASL dataset.

Model training:

```
for epoch in range(num_epochs):
    # Training
    model.train()
    total_train = 0
    correct_train = 0
    train_loss = 0

    for images, labels in train:
        batch_size = images.size(0)  # Get current batch size dynamically
        images = Variable(images.view(batch_size, 1, 32, 32))
        labels = Variable(labels)
        optimizer.zero_grad()
        outputs = model(images)
        loss = error(outputs, labels)
        loss.backward()
        optimizer.step()
        train_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total_train += labels.size(0)
        correct_train += (predicted == labels).sum().item()

    train_accuracy = 100 * correct_train / total_train
    avg_train_loss = train_loss / len(train)

    # Validation
    model.eval()
    total_val = 0
    correct_val = 0
    val_loss = 0

    with torch.no_grad():
        for images, labels in test:
            batch_size = images.size(0)  # Get current batch size dynamically
            images = Variable(images.view(batch_size, 1, 32, 32))
            labels = Variable(labels)
            outputs = model(images)
            loss = error(outputs, labels)
            val_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total_val += labels.size(0)
            correct_val += (predicted == labels).sum().item()

    val_accuracy = 100 * correct_val / total_val
    avg_val_loss = val_loss / len(test)

    # Print training and validation results
    print(f'Epoch {epoch + 1}, Loss: {avg_train_loss:.4f}, Accuracy: {train_accuracy:.2f}%, Val Loss: {avg_val_loss:.4f}, Val Accuracy: {val_accuracy:.2f}%')
```

Figure 33 (train Pytorch CNN model)

The code details the training and validation process for a convolutional neural network (CNN) over multiple epochs. For each epoch, the model is set to training mode (model.train()), and metrics for total training loss, total training examples, and correct predictions are initialized. Within the training loop, the images and labels are reshaped, and the model makes predictions. The loss is computed and backpropagated, and the optimizer updates the model's parameters. The training accuracy and average training loss are calculated after iterating through all batches. For validation, the model is set to evaluation mode (model.eval()), and gradients are disabled (torch.no_grad()). The validation loop processes the validation data, similarly, computing the validation loss and accuracy without updating the model parameters. After each epoch, the training and validation metrics are printed. This comprehensive approach ensures that the model is trained effectively while monitoring its performance on both the training and validation datasets.

10.2.2 ARSL Model:

consist of CNN is the final model with high accuracy.

10.2.2.1 CNN:

Preprocessing:

```
# global variables
Language = "Ar"
ImageClassMapping_path = "../input/Labels/ImagesClassPath.csv"
ClassLabels_path = "../input/Labels/ClassLabels.xlsx"
ImagesRoot_path = "../input/"

ModelFileName ='Model_255.h5'
```

Figure 34 (define the data and language)

The code defines several global variables used for configuring the dataset and model paths. Language is set to "Ar", indicating the language setting. ImageClassMapping_path specifies the path to a CSV file containing the mapping of images to their respective classes. ClassLabels_path points to an Excel file that contains the class labels. ImagesRoot_path sets the root directory where the input images are stored. Lastly, ModelFileName is assigned the name "Model_255.h5", which is likely the filename for saving the trained model. These variables centralize important file paths and settings, making the code more organized and easier to maintain.

```
# load 54k image path mapping
df_ImageClassPath = pd.read_csv(ImageClassMapping_path)
display(df_ImageClassPath.head())
```

Figure 35 (map creation)

The code loads a CSV file containing image path mappings into a pandas DataFrame. The CSV file path is specified by the variable ImageClassMapping_path. The pd.read_csv function reads the CSV file and stores the data in the DataFrame df_ImageClassPath. The display function is then used to show the first few rows of the DataFrame with

df_ImageClassPath.head(), providing a quick view of the data structure and contents. This step is essential for verifying that the image paths and their corresponding class mappings have been loaded correctly, facilitating further data processing and analysis.

```
# load Class Labels
df_Classes = pd.read_excel(ClassLabels_path)
display(df_Classes.head())
```

Figure 36 (label map)

The code loads an Excel file containing class labels into a pandas DataFrame. The Excel file path is specified by the variable ClassLabels_path. The pd.read_excel function reads the Excel file and stores the data in the DataFrame df_Classes. The display function is then used to show the first few rows of the DataFrame with df_Classes.head(), providing a quick view of the data structure and contents. This step is essential for verifying that the class labels have been loaded correctly, facilitating further data processing and analysis for tasks such as classification.

```
df_ImageClassPath.groupby("ClassId").size().describe()
```

Figure 37 (view number of classes)

The code groups the DataFrame df_ImageClassPath by the column "ClassId" and calculates the size of each group. The groupby method is used to group the data by "ClassId", and the size method counts the number of occurrences in each group. The describe method is then called on the resulting series to generate descriptive statistics, such as the count, mean, standard deviation, minimum, maximum, and quartiles, for the sizes of the groups. This step provides a summary of the distribution of image counts across different classes, helping to understand the class balance in the dataset.

```
ddata = {"samples destribution":df_ImageClassPath.groupby("ClassId").size()}
iindex = range(32)

ddataframe = pd.DataFrame(data=ddata, index= iindex)
ddataframe.plot.bar(stacked= True, rot= 15, title='samples destribution')
plt.show(block= True)
```

Figure 38 (print sample with its label)

The code creates a bar plot to visualize the distribution of samples across different classes in the dataset. It first groups the df_ImageClassPath DataFrame by "ClassId" and calculates the size of each group, storing the result in a dictionary ddata with the key "samples distribution". An index iindex is defined as a range of 32, representing the class indices. A new DataFrame ddataframe is created from the ddata dictionary using the defined index. The plot.bar method is then used to create a stacked bar plot with the title "samples distribution". The bars are rotated by 15 degrees for better readability, and

plt.show(block=True) is called to display the plot. This visualization helps in understanding the distribution and balance of samples across different classes in the dataset.

```
def SplitData(predictions,testsize):

    min = df_ImageClassPath.groupby("ClassId").size().min()
    print('{0} Samples per Class'.format(min))

    # empty dataframes with same column difinition
    df_TrainingSet = df_ImageClassPath[0:0].copy()
    df_TestSet = df_ImageClassPath[0:0].copy()
    df_PredSet = df_ImageClassPath[0:0].copy()

    # Create the sets by loop thru classes and append
    for index,row in df_Classes.iterrows():
        # make sure all class are same size
        df_FullSet = df_ImageClassPath[df_ImageClassPath['ClassId'] == row['ClassId']].sample(min,random_state= 42)

#         df_FullSet = df_ImageClassPath[df_ImageClassPath['ClassId'] == row['ClassId']]

        df_PredSet = df_PredSet.append(df_FullSet.sample(n=predictions, random_state=1))
        df_FullSet = pd.merge(df_FullSet,df_PredSet, indicator=True,
                              how='left').query('_merge=="left_only"').drop('_merge', axis=1)

        trainingSet, testSet = train_test_split(df_FullSet, test_size= testsize)

        df_TrainingSet = df_TrainingSet.append(trainingSet)
        df_TestSet = df_TestSet.append(testSet)

    return df_TrainingSet,df_TestSet,df_PredSet
```

Figure 39 (split the data)

The code defines a function SplitData that splits a dataset into training, testing, and prediction sets based on class distributions. It begins by determining the minimum number of samples per class from df_ImageClassPath and prints this value. Empty DataFrames (df_TrainingSet, df_TestSet, and df_PredSet) with the same structure as df_ImageClassPath are initialized. The function then iterates through each class in df_Classes, ensuring an equal number of samples for each class by sampling min samples per class. During each iteration, it updates df_PredSet with a random subset of samples for predictions, removes these from the full set, and splits the remaining data into training and testing sets using train_test_split with the specified testsize. The resulting training, testing, and prediction sets are appended to their respective DataFrames. Finally, the function returns df_TrainingSet, df_TestSet, and df_PredSet. This process ensures balanced datasets for training and evaluation, and a separate set for making predictions.

```python
def get_classlabel(class_code,lang= 'Ar'):
    if lang== 'Ar':
        text_to_be_reshaped = df_Classes.loc[df_Classes['ClassId'] == class_code,
                                             'ClassAr'].values[0]
        reshaped_text = arabic_reshaper.reshape(text_to_be_reshaped)
        return get_display(reshaped_text)
    elif lang== 'En':
        return df_Classes.loc[df_Classes['ClassId'] == class_code, 'Class'].values[0]
```

Figure 40 (retrieve class label)

The code defines a function get_classlabel that retrieves class labels based on a given class code and language preference. The function takes two arguments: class_code and lang (with a default value of 'Ar'). If the language is Arabic ('Ar'), the function extracts the Arabic class label from df_Classes where ClassId matches class_code, reshapes the text for proper display using arabic_reshaper, and returns the reshaped text formatted with get_display. If the language is English ('En'), the function directly returns the English class label from df_Classes where ClassId matches class_code. This function facilitates multilingual support for retrieving class labels, ensuring proper display of Arabic text.

```python
def getDataSet(setType,isDL): # 'Training' for Training dataset , 'Testing' for Testing data set
    imgs = []
    lbls = []
    df = pd.DataFrame(None)

    if setType =='Training':
        df = dtTraining.copy()
    elif setType=='Test':
        df = dtTest.copy()
    elif setType=='Prediction':
        df = dtPred.copy()

    for index,row in df.iterrows():
        lbls.append(row['ClassId'])
        try:
            imageFilePath = os.path.join(ImagesRoot_path, row['ImagePath'])
            img = image.load_img(imageFilePath, target_size=(64,64,1),
                                 color_mode = "grayscale")
            img = image.img_to_array(img) # to array
            img = img/255 # Normalize
            if isDL == False:
                img = img.flatten() # for knn_classifier Model
            imgs.append(img)

        except Exception as e:
            print(e)

    shuffle(imgs,lbls,random_state=255) #Shuffle the dataset

    imgs = np.array(imgs)
    lbls = np.array(lbls)
    if isDL ==True:
        lbls = to_categorical(lbls) # for keras CNN Model
    return imgs, lbls
```

Figure 41 (load the data)

function getDataSet that prepares image data and labels for training, testing, or prediction. The function takes two parameters: setType, which specifies whether to prepare 'Training', 'Testing', or 'Prediction' data, and isDL, a boolean indicating if the data is for deep learning models. Based on setType, the corresponding DataFrame (dfTraining, dfTest, or dfPred) is copied to df. The function initializes empty lists for images (imgs) and labels (lbls). It iterates over each row in the DataFrame, extracting the class ID and image path. The images are loaded, resized to 64x64 pixels, converted to grayscale, and normalized by dividing by 255. If isDL is False, the images are flattened for use with KNN classifiers; otherwise, they remain as arrays for CNN models. The images and labels are then shuffled and converted to NumPy arrays. If isDL is True, the labels are converted to categorical format suitable for Keras CNN models. The function returns the processed images and labels, facilitating the preparation of datasets for the task.

```
def display_prediction(col_size, row_size,XPred,yPred):
    img_index=0
    fig, ax = plt.subplots(row_size, col_size, figsize=(row_size*2.5,col_size*1.5))
    for row in range(0,row_size):
        for col in range(0,col_size):
            ax[row][col].imshow(XPred[img_index][:,:,0], cmap='gray')
            ax[row][col].set_title("({}) {}".format(yPred[img_index],get_classlabel(yPred[img_index],'Ar')))
            ax[row][col].set_xticks([])
            ax[row][col].set_yticks([])
            img_index += 1
```

Figure 42 (display prediction)

The code defines a function display_prediction that visualizes a grid of predicted images along with their corresponding class labels. The function takes four parameters: col_size and row_size to define the grid dimensions, and XPred and yPred for the predicted images and labels, respectively. It initializes an image index (img_index) and creates a figure with subplots using plt.subplots, setting the figure size based on the grid dimensions. The function then iterates through each subplot cell, displaying each image from XPred in grayscale using imshow. The subplot title is set to the predicted class label obtained by calling get_classlabel with the current label from yPred. The x-ticks and y-ticks are removed for a cleaner presentation. The image index is incremented in each iteration to move to the next image. This function provides a visual summary of the model's predictions, allowing for an easy assessment of its performance.

```
ddata = {"Training":dtTraining.groupby("ClassId").size(),"Test":dtTest.groupby("ClassId").size()}
iindex = range(32)

ddataframe = pd.DataFrame(data=ddata, index= iindex)
ddataframe.plot.bar(stacked= True, rot= 15, title='Training vs Test data')
plt.show(block= True)
```

Figure 43 (group train and test data by id)

The code creates a bar plot to visualize and compare the distribution of samples across different classes in the training and testing datasets. It constructs a dictionary ddata that includes the sizes of each class in the dtTraining and dtTest DataFrames, grouped by ClassId. An index iindex is defined as a range of 32, representing the class indices. This dictionary is

converted into a pandas DataFrame ddataframe using the defined index. The plot.bar method is then called on the DataFrame to create a stacked bar plot, with the bars rotated by 15 degrees for better readability. The plot is titled "Training vs Test data", and plt.show(block=True) is used to display the plot. This visualization helps in assessing the balance and distribution of the training and testing datasets across different classes.

Model creating:

```
X_train,y_train = getDataSet('Training',True)
X_valid,y_valid = getDataSet('Test',True)
X_pred,_ = getDataSet('Prediction',True)
```

Figure 44 (train test preparation)

The code uses the getDataSet function to prepare datasets for training, validation, and prediction. The function is called with the setType parameter set to 'Training', 'Test', and 'Prediction', and the isDL parameter set to True, indicating the data is intended for deep learning models. The resulting datasets are assigned to variables X_train and y_train for the training set, X_valid and y_valid for the validation set, and X_pred for the prediction set. This setup ensures that the data is appropriately processed and ready for use in training, validating, and making predictions with a deep learning model.

```
model = Models.Sequential()

model.add(Layers.Conv2D(64, kernel_size=(3, 3),activation='relu',input_shape=(64,64,1)))
model.add(Layers.Conv2D(64, (3, 3), activation='relu'))
model.add(Layers.MaxPooling2D(pool_size=(2, 2)))
model.add(Layers.Dropout(0.25))
model.add(Layers.Flatten())
model.add(Layers.Dense(128, activation='relu'))
model.add(Layers.Dropout(0.5))
model.add(Layers.Dense(32, activation='softmax'))

model.compile(loss='categorical_crossentropy',optimizer='Adam',metrics=['accuracy'])

model.summary()
Utils.plot_model(model,to_file='model.png',show_shapes=True, show_layer_names=True, dpi=80)
```

Figure 45 (model layers)

The code defines and compiles a convolutional neural network (CNN) model using Keras' Sequential API. The model starts with two convolutional layers (Conv2D), each with 64 filters, a kernel size of 3x3, and ReLU activation. The first convolutional layer includes an input_shape of 64x64x1 for grayscale images. A max pooling layer (MaxPooling2D) with a pool size of 2x2 follows, along with a dropout layer (Dropout) with a dropout rate of 0.25 to prevent overfitting. The model then flattens the feature maps using Flatten and adds a fully connected dense layer (Dense) with 128 units and ReLU activation, followed by another dropout layer with a dropout rate of 0.5. The final dense layer has 32 units with softmax activation, suitable for multi-class classification. The model is compiled with categorical cross-entropy loss, the Adam optimizer, and accuracy as the evaluation metric. A summary

of the model architecture is printed, and a visual representation of the model is saved to 'model.png' using plot_model, showing shapes and layer names with a DPI of 80.

Model training:

```
callbacks_list =[EarlyStopping(monitor='val_loss', patience=10), ModelCheckpoint(
    filepath='model_255.h5', monitor='val_loss', save_best_only= True),]

trained = model.fit(X_train, y_train, epochs=35, validation_data=(X_valid, y_valid),
                    callbacks= callbacks_list)
```

Figure 46 (early stopping function and model training)

The code sets up callbacks for training a Keras model and then fits the model to the training data. The callbacks_list includes two callbacks: EarlyStopping, which monitors the validation loss (val_loss) and stops training if it doesn't improve for 10 epochs (patience=10), and ModelCheckpoint, which saves the best model to 'model_255.h5' based on the lowest validation loss, ensuring only the best version of the model is saved (save_best_only=True). The model.fit method is called to train the model on the training data (X_train, y_train) for 35 epochs, using the validation data (X_valid, y_valid) to monitor performance. The callbacks_list is passed to the fit method to apply the specified callbacks during training. This setup helps in preventing overfitting and saving the best model based on validation performance.

10.2.3 ASL Realtime:

10.2.3.1 SVM:

```
# Function to preprocess the image and extract HOG features
def preprocess_image(image):
    try:
        image = transform.resize(image, (64, 64))  # Resize to 64x64
        image = color.rgb2gray(image)  # Convert to grayscale
        hog_features = feature.hog(image, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
        return hog_features
    except Exception as e:
        print(f"Error in preprocessing image: {e}")
        return None
```

Figure 47 (data preprocessing)

The code defines a function preprocess_image that preprocesses an image and extracts Histogram of Oriented Gradients (HOG) features. The function takes an image as input and performs the following steps within a try block: it resizes the image to 64x64 pixels using transform.resize, converts the image to grayscale using color.rgb2gray, and then extracts HOG features using the feature.hog function with specified parameters for pixels per cell and cells per block. The HOG features are returned. If an exception occurs during these steps, an error message is printed, and the function returns None. This function is essential for preparing image data by normalizing its size and extracting relevant features for further processing in machine learning tasks.

```
# Initialize MediaPipe Hands
mp_hands = mp.solutions.hands
hands = mp_hands.Hands(static_image_mode=False,
                       max_num_hands=1,
                       min_detection_confidence=0.5,
                       min_tracking_confidence=0.5)
```

Figure 48 (media pipe hand initialization)

The code initializes the MediaPipe Hands solution for hand detection and tracking. The mp_hands module from mp.solutions is used to access the Hands API. An instance of the Hands class is created and assigned to the variable hands with the following parameters: static_image_mode=False indicates that the model will treat the input images as a video stream, allowing for continuous detection; max_num_hands=1 restricts the detection to a single hand; min_detection_confidence=0.5 sets the minimum confidence threshold for the hand detection model to consider a detection as valid; and min_tracking_confidence=0.5 sets the minimum confidence threshold for the hand tracking model to consider the tracked landmarks as reliable. This setup prepares the MediaPipe Hands solution for real-time hand detection and tracking, useful in applications such as gesture recognition and sign language interpretation.

```python
# Initialize the webcam
cap = cv2.VideoCapture(0)
if not cap.isOpened():
    print("Error: Could not open webcam.")
    exit()

try:
    while True:
        ret, frame = cap.read()
        if not ret:
            print("Failed to capture frame. Exiting...")
            break

        # Convert the frame to RGB as MediaPipe expects RGB images
        frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)

        # Use MediaPipe to detect hands in the frame
        results = hands.process(frame_rgb)

        if results.multi_hand_landmarks:
            # A hand is detected, proceed with preprocessing and prediction
            hog_features = preprocess_image(frame_rgb)

            if hog_features is not None:
                print(f"HOG Features: {hog_features[:10]}")  # Print first 10 HOG features for inspection
                # Reshape for the model (1 sample, number of features)
                hog_features = np.array(hog_features).reshape(1, -1)
                print(f"HOG Features Shape: {hog_features.shape}")

                # Perform inference
                try:
                    predicted_index = model.predict(hog_features)[0]
                    print(f"Predicted Index: {predicted_index}")
                except Exception as e:
                    print(f"Error during prediction: {e}")
                    predicted_index = '?'

                # Draw hand landmarks on the frame
                for hand_landmarks in results.multi_hand_landmarks:
                    mp.solutions.drawing_utils.draw_landmarks(frame, hand_landmarks, mp_hands.HAND_CONNECTIONS)

                # Display the prediction on the frame
                cv2.putText(frame, f'Predicted Index: {predicted_index}', (10, 50), cv2.FONT_HERSHEY_SIMPLEX, 1, (0, 255, 0), 2, cv2.LINE_AA)
            else:
                print("HOG features extraction failed.")
                cv2.putText(frame, 'Error in preprocessing', (10, 50), cv2.FONT_HERSHEY_SIMPLEX, 1, (0, 0, 255), 2, cv2.LINE_AA)
        else:
            cv2.putText(frame, 'No hand detected', (10, 50), cv2.FONT_HERSHEY_SIMPLEX, 1, (0, 0, 255), 2, cv2.LINE_AA)

        # Display the resulting frame
        cv2.imshow('Real-Time Sign Language Recognition', frame)

        if cv2.waitKey(1) & 0xFF == ord('q'):
            break
finally:
    cap.release()
    cv2.destroyAllWindows()
    hands.close()
```

Figure 49 (real time prediction)

The code sets up a real-time sign language recognition system using OpenCV and MediaPipe. It initializes the webcam for video capture and enters a loop to process each frame. The frame is converted from BGR to RGB, as required by MediaPipe, and passed to the hands.process method to detect hand landmarks. If hand landmarks are detected, the image is preprocessed to extract Histogram of Oriented Gradients (HOG) features using the preprocess_image function. These features are reshaped and used to predict the sign language gesture with a pre-trained model. The predicted index is displayed on the frame. If any errors occur during preprocessing or prediction, appropriate messages are displayed. Hand landmarks are drawn on the frame, and the annotated frame is shown in a window titled "Real-Time Sign Language Recognition". The loop continues until the 'q' key is pressed.

Upon exiting, the webcam is released, and all OpenCV windows are closed. This system allows for real-time detection and recognition of sign language gestures.

10.2.3.2 Efficient net:

```
# Load the EfficientNet model with the custom number of output classes
model_name = 'efficientnet-b0'  # Adjust the model variant if needed
model = EfficientNet.from_name(model_name, num_classes=29)

# Specify the path to your model's state dict
model_path = r"E:\gp_final\image\recognition\EfficientNet\ASL_HandSignLang_EfficientNetB0_2.pth"

# Load the state dict
state_dict = torch.load(model_path, map_location=torch.device('cpu'))

# Remove the 'module.' prefix if it exists
new_state_dict = {}
for k, v in state_dict.items():
    if k.startswith('module.'):
        new_state_dict[k[7:]] = v
    else:
        new_state_dict[k] = v

# Load the cleaned state dict into the model
model.load_state_dict(new_state_dict)

# Now the model should be loaded correctly
print("Model loaded successfully!")
```

Figure 50 (load trained model)

The code loads a pre-trained EfficientNet model with a custom number of output classes and adjusts the model state dictionary to ensure compatibility. First, the EfficientNet model is instantiated with the specified variant (efficientnet-b0) and set to have 29 output classes. The path to the model's state dictionary file is defined by model_path, and the state dictionary is loaded using torch.load, with mapping to the CPU device. A new state dictionary, new_state_dict, is created by iterating over the original state dictionary and removing any 'module.' prefix from the keys. This adjustment ensures compatibility, especially when the model was trained using DataParallel. The cleaned state dictionary is then loaded into the model using model.load_state_dict. Finally, a confirmation message "Model loaded successfully!" is printed, indicating that the model has been correctly loaded and is ready for use.

```
model.eval()  # Set the model to evaluation mode

# Setup MediaPipe Hands
mp_hands = mp.solutions.hands
hands = mp_hands.Hands(static_image_mode=False,
                       max_num_hands=1,
                       min_detection_confidence=0.5,
                       min_tracking_confidence=0.5)
```

Figure 51 (media pipe hand setup)

The code sets up the model and the MediaPipe Hands solution for evaluation. First, the model.eval() method is called to set the model to evaluation mode, which disables certain layers such as dropout and batch normalization, making the model behavior consistent during inference. Then, the MediaPipe Hands solution is initialized by creating an instance of mp_hands.Hands with specified parameters: static_image_mode=False for treating the input as a video stream, max_num_hands=1 to detect only one hand, min_detection_confidence=0.5 to set the threshold for considering a hand detection valid, and min_tracking_confidence=0.5 for the threshold of reliable tracking. This setup prepares both the model and the hand detection system for accurate and stable performance during real-time sign language recognition tasks.

```
# Define the transformation
transform = transforms.Compose([
    transforms.Resize((128, 128)),  # Resize the image to 128x128
    transforms.ToTensor()           # Convert the image to a PyTorch tensor
])
```

Figure 52 (images resize)

The code defines a series of transformations to be applied to images using PyTorch's transforms module. These transformations are composed using transforms.Compose, which chains multiple transformations together. The first transformation, transforms.Resize(128, 128), resizes the input image to 128x128 pixels. The second transformation, transforms.ToTensor(), converts the resized image into a PyTorch tensor. This setup ensures that all images processed through this transformation pipeline will be uniformly resized and converted into a format suitable for input model, facilitating consistent preprocessing for model training or inference.

```python
cap = cv2.VideoCapture(0)
if not cap.isOpened():
    print("Error: Could not open webcam.")
    exit()

# Define label mapping
label_map = {
    0: 'A', 1: 'B', 2: 'C', 3: 'D', 4: 'E', 5: 'F', 6: 'G', 7: 'H', 8: 'I', 9: 'J',
    10: 'K', 11: 'L', 12: 'M', 13: 'N', 14: 'O', 15: 'P', 16: 'Q', 17: 'R', 18: 'S', 19: 'T',
    20: 'U', 21: 'V', 22: 'W', 23: 'X', 24: 'Y', 25: 'Z', 26: 'del', 27: 'nothing', 28: 'space'
}

# Variables to track the prediction, timing, and word construction
last_predicted_label = None
label_change_time = time.time()
consistency_duration = 3  # Duration for which the prediction must remain the same
word_buffer = []
last_hand_detected_time = time.time()
hand_detected = False

try:
    while True:
        ret, frame = cap.read()
        if not ret:
            print("Failed to capture frame. Exiting...")
            break

        # Convert the frame to RGB and then to a PIL Image
        image = Image.fromarray(cv2.cvtColor(frame, cv2.COLOR_BGR2RGB))

        # Apply the transformations
        input_tensor = transform(image)
        input_batch = input_tensor.unsqueeze(0)  # Create a mini-batch as expected by the model

        # Perform inference
        with torch.no_grad():
            output = model(input_batch)
            predicted_index = output.argmax(1).item()  # Get the index of the max log-probability
            predicted_label = label_map.get(predicted_index, '?')  # Map index to label

        # Check if the predicted label has changed
        if predicted_label != last_predicted_label:
            last_predicted_label = predicted_label
            label_change_time = time.time()
        elif time.time() - label_change_time >= consistency_duration:
            print(f"Confirmed Prediction: {predicted_label}")
            if predicted_label != 'nothing':
                word_buffer.append(predicted_label)
            last_predicted_label = None  # Reset after printing to prevent continuous printing

        # Use MediaPipe to draw hand landmarks on the frame
        frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
        results = hands.process(frame_rgb)
        if results.multi_hand_landmarks:
            hand_detected = True
            last_hand_detected_time = time.time()
            for hand_landmarks in results.multi_hand_landmarks:
                mp.solutions.drawing_utils.draw_landmarks(frame, hand_landmarks, mp_hands.HAND_CONNECTIONS)
        else:
            hand_detected = False

        # Check if no hand detected for 5 seconds to end the word
        if not hand_detected and time.time() - last_hand_detected_time >= 5:
            if word_buffer:
                print(f"Word: {''.join(word_buffer)}")
                word_buffer = []

        cv2.putText(frame, f'Predicted Label: {predicted_label}', (10, 50), cv2.FONT_HERSHEY_SIMPLEX, 1, (0, 255, 0), 2, cv2.LINE_AA)
        cv2.putText(frame, f'Current Word: {"".join(word_buffer)}', (10, 100), cv2.FONT_HERSHEY_SIMPLEX, 1, (255, 0, 0), 2, cv2.LINE_AA)
        cv2.imshow('Real-Time Sign Language Recognition', frame)

        if cv2.waitKey(1) & 0xFF == ord('q'):
            break
finally:
    cap.release()
    cv2.destroyAllWindows()
    hands.close()
```

Figure 53 (real time prediction)

The code implements a real-time sign language recognition system using OpenCV, PyTorch, and MediaPipe. The webcam is initialized using cv2.VideoCapture(0), and a dictionary label_map is defined to map prediction indices to corresponding sign language letters and

actions. Variables are initialized to track the last predicted label, timing, and word construction. The main loop reads frames from the webcam, converts them to RGB, and applies the predefined transformations to resize and convert the image to a PyTorch tensor. The model predicts the label for the current frame, and the predicted index is mapped to a label using label_map.

If the predicted label changes, the label and the time of change are updated. If the label remains consistent for a defined duration, it is appended to word_buffer for word construction, except for 'nothing'. MediaPipe is used to detect and draw hand landmarks on the frame. If no hand is detected for 5 seconds, the constructed word from word_buffer is printed, and the buffer is reset.

The frame is annotated with the predicted label and the current constructed word, then displayed in a window titled "Real-Time Sign Language Recognition". The loop continues until the 'q' key is pressed. Upon exit, the webcam is released, all OpenCV windows are closed, and the MediaPipe Hands instance is closed, ensuring proper cleanup of resources.

10.2.4 ASL word spelling correction:

I have trained this model on 90K words.

```
# Load the Word2Vec model
model = KeyedVectors.load(r"E:\gp_final\image\word correction\word2vec\word2vec_model.bin")

# Load the vocabulary
with open(r"E:\gp_final\image\word correction\word2vec\vocabulary.pkl", 'rb') as f:
    vocabulary = pickle.load(f)

# Function to find the nearest word based on Levenshtein distance
def find_nearest_word(user_word, word_list):
    nearest_word = None
    min_distance = float('inf')

    for word in word_list:
        distance = Levenshtein.distance(user_word, word)
        if distance < min_distance:
            min_distance = distance
            nearest_word = word

    return nearest_word

# Example usage
user_word = input("Enter a word: ")
nearest_word = find_nearest_word(user_word, vocabulary)
print(f"The word nearest to '{user_word}' is '{nearest_word}'")
```

Figure 54 (word spelling correction)

The code loads a pre-trained Word2Vec model and a vocabulary list, and defines a function to find the nearest word based on Levenshtein distance. The Word2Vec model is loaded from a binary file using KeyedVectors.load, and the vocabulary list is loaded from a pickle file. The function find_nearest_word takes a user_word and a word_list, calculates the Levenshtein distance between user_word and each word in the word_list, and returns the word with the smallest distance. An example usage is provided where the user is prompted to input a word, and the nearest word from the vocabulary is found and printed. This setup

allows for correcting real-time predicted words by finding the closest match from a large vocabulary using both Word2Vec embeddings and Levenshtein distance.

10.3 Video recognition system:

I have searched and try to collect dataset for this part I also try to create one but there is good data set for this task, so I used dataset for Brazilian sign language consist of frames of the videos and this dataset consists of 3066 frames.(there are like 10 more trials with another models LSTM, CNN, and RNN but all of this with different datasets that was very bad).

10.3.1 BSL

10.3.1.1 VGG19

Preprocessing:

```
# Define the paths to the dataset
data_path = '/kaggle/input/brazilian-sign-language-words-recognition/30 FPS/30 FPS/train'
# Function to load images and labels
def load_data(data_path):
    data = []
    labels = []
    for word_folder in os.listdir(data_path):
        if word_folder.lower() == 'formation':  # Ignore the 'formation' class
            continue
        word_folder_path = os.path.join(data_path, word_folder)
        if os.path.isdir(word_folder_path):
            for frame_file in os.listdir(word_folder_path):
                frame_path = os.path.join(word_folder_path, frame_file)
                img = cv2.imread(frame_path)
                if img is not None:
                    img = cv2.resize(img, (64, 64))  # Resize all images to 224x224 for VGG19
                    data.append(img)
                    labels.append(word_folder)
    return np.array(data), np.array(labels)
```

Figure 55 (load data)

The code defines a function load_data that loads images and their corresponding labels from a specified dataset path for training a model. The data_path variable specifies the root directory containing the dataset. The load_data function initializes two empty lists, data and labels, to store images and labels, respectively. It iterates through each folder in the dataset directory, ignoring any folder named 'formation'. For each valid word folder, it constructs the path and iterates through the image files within that folder. Each image is read using OpenCV's cv2.imread, resized to 64x64 pixels, and appended to the data list. The corresponding label (folder name) is appended to the labels list. Finally, the function returns the images and labels as NumPy arrays. This function facilitates the preprocessing of image data, making it ready for training machine learning models.

```
# Load data
X, y = load_data(data_path)

# Normalize image data
X = X / 255.0

# Encode labels
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
```

Figure 56 (label encoder and normalization)

The code demonstrates the process of loading and preprocessing image data for a machine learning model. First, the load_data function is called with data_path as an argument, returning the images and their corresponding labels stored in X and y, respectively. The image data in X is then normalized by dividing each pixel value by 255.0 to scale the values between 0 and 1, which is a common preprocessing step for image data. Next, a LabelEncoder instance is created to encode the labels. The fit_transform mtraining andlEncoder is applied to y, transforming the categorical labels into a numeric format suitable for model training, and storing the encoded labels in y_encoded. This snippet effectively prepares the image data and labels for training a machine learning model by normalizing the images and encoding the labels.

```
# Data Augmentation
datagen = ImageDataGenerator(
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True,
    fill_mode='nearest'
)

# Augment data to ensure each class has 1000 samples
augmented_data = []
augmented_labels = []
```

Figure 57 (data augmentation)

The code sets up data augmentation for image preprocessing using Keras' ImageDataGenerator. The datagen instance is configured with several augmentation parameters: rotation_range of 10 degrees, width_shift_range and height_shift_range of 0.1, shear_range of 0.1, zoom_range of 0.1, and horizontal_flip set to True to randomly flip images horizontally. The fill_mode is set to 'nearest' to fill in new pixels that may appear after transformations. These augmentations help in creating a more robust model by artificially increasing the diversity of the training dataset. Additionally, two empty lists, augmented_data and augmented_labels, are initialized to store the augmented images and

their corresponding labels. This setup aims to ensure that each class in the dataset has 1,000 samples through data augmentation, improving the model's generalization capability.

```
for label in np.unique(y_encoded):
    class_data = X[y_encoded == label]
    class_labels = y_encoded[y_encoded == label]
    if len(class_data) < 1000:
        augment_iter = datagen.flow(class_data, class_labels, batch_size=1)
        while len(class_data) < 1000:
            augmented_image, augmented_label = next(augment_iter)
            class_data = np.append(class_data, augmented_image, axis=0)
            class_labels = np.append(class_labels, augmented_label, axis=0)
    augmented_data.extend(class_data[:1000])
    augmented_labels.extend(class_labels[:1000])

X_augmented = np.array(augmented_data)
y_augmented = np.array(augmented_labels)
```

Figure 58 (save data in arrays)

The code augments the dataset to ensure that each class has at least 1,000 samples using Keras' ImageDataGenerator. It iterates over each unique label in y_encoded. For each label, it selects the corresponding images (class_data) and labels (class_labels). If the number of samples for a particular class is less than 1,000, it uses datagen.flow to create an iterator that generates augmented images and labels. In a while loop, it continuously augments and appends images to class_data and labels to class_labels until the class has 1,000 samples. The augmented data and labels are then added to augmented_data and augmented_labels lists, respectively. After processing all classes, augmented_data and augmented_labels are converted to NumPy arrays, X_augmented and y_augmented. This ensures that the dataset is balanced with a minimum of 1,000 samples per class, improving the model's ability to generalize across all classes.

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_augmented, y_augmented, test_size=0.2, random_state=42)
```

Figure 59 (train and test split)

The code splits the augmented dataset into training and testing sets using the train_test_split function from the sklearn.model_selection module. The augmented data and labels, X_augmented and y_augmented, are divided into training and testing subsets. The test_size=0.2 parameter specifies that 20% of the data should be allocated for testing, while the remaining 80% will be used for training. The random_state=42 ensures that the split is reproducible by initializing the random number generator to a fixed state. The resulting subsets are stored in X_train, X_test, y_train, and y_test, providing separate datasets for model training and evaluation. This step is crucial for assessing the model's performance on unseen data and ensuring it generalizes well.

<u>Model training:</u>

```python
# Load the VGG19 model with pre-trained ImageNet weights, excluding the top layer
base_model = VGG19(weights='imagenet', include_top=False, input_shape=(64,64,3))

# Add custom top layers
x = base_model.output
x = Flatten()(x)
x = Dense(512, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(256, activation='relu')(x)
x = Dropout(0.3)(x)
predictions = Dense(len(np.unique(y_train)), activation='softmax')(x)

# Define the full model
model = Model(inputs=base_model.input, outputs=predictions)

# Freeze the base VGG19 layers
for layer in base_model.layers:
    layer.trainable = False

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

Figure 60 (model layers)

The code builds a deep learning model by leveraging the pre-trained VGG19 architecture and customizing the top layers for a specific classification task. First, the VGG19 model is loaded with pre-trained ImageNet weights, excluding the top layer (include_top=False), and the input shape is set to (64, 64, 3). Custom top layers are then added: the output of the VGG19 base model is flattened, followed by a dense layer with 512 units and ReLU activation, a dropout layer with a 20% dropout rate, another dense layer with 256 units and ReLU activation, and another dropout layer with a 30% dropout rate. The final output layer has a number of units equal to the number of unique classes in y_train and uses softmax activation for classification. The full model is defined with the input from the VGG19 base model and the custom top layers as output. The base VGG19 layers are frozen to prevent them from being updated during training. Finally, the model is compiled with the Adam optimizer, sparse categorical cross-entropy loss, and accuracy as the evaluation metric, preparing it for training on the specific dataset.

```python
# Train the model
history = model.fit(X_train, y_train, epochs=15, validation_data=(X_test, y_test))
```

Figure 61 (model training)

The code initiates the training process for the previously defined model using the fit method. The model is trained on the X_train and y_train datasets for 15 epochs. During each epoch, the model's performance is validated using the X_test and y_test datasets. The results of the training and validation processes are stored in the history object, which can be used later to analyze the model's learning curves and performance metrics. This step is crucial for adjusting the model's parameters to minimize the loss function and improve its accuracy on both the training and validation data.

```python
# Unfreeze the base VGG19 layers and re-train
for layer in base_model.layers:
    layer.trainable = True

# Compile the model again with a lower learning rate
model.compile(optimizer=tf.keras.optimizers.Adam(1e-5), loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Continue training
history_fine = model.fit(X_train, y_train, epochs=15, validation_data=(X_test, y_test))
```

Figure 62 (train with base layers)

The code continues the model training process by first unfreezing the previously frozen base VGG19 layers to allow fine-tuning. Each layer in the base model is set to trainable = True, enabling their weights to be updated during training. The model is then recompiled with a significantly lower learning rate (1e-5) using the Adam optimizer, to facilitate fine-tuning without overshooting the minima in the loss landscape. The loss function remains as sparse categorical cross-entropy, and accuracy is retained as the metric. The training process is continued by fitting the model on the X_train and y_train datasets for another 15 epochs, with X_test and y_test used for validation. The results are stored in the history_fine object, which allows for comparison with the initial training phase and assessment of the improvements brought by fine-tuning.

10.4 Video Request:

```python
# Define the path to your dataset
dataset_path = r"E:\gp_final\word to video\videos_by_gloss"

def list_subfolders(dataset_path):
    subfolders = [f for f in os.listdir(dataset_path) if os.path.isdir(os.path.join(dataset_path, f))]
    return subfolders

def play_video(video_path):
    cap = cv2.VideoCapture(video_path)

    if not cap.isOpened():
        print(f"Error: Cannot open video {video_path}")
        return

    # Get the frame rate of the video
    fps = cap.get(cv2.CAP_PROP_FPS)
    if fps == 0:
        print("Error: Unable to determine frame rate of the video")
        return

    delay = int(1000 / fps)  # Calculate delay in milliseconds

    while cap.isOpened():
        ret, frame = cap.read()
        if not ret:
            break

        # Display the frame in a popup window
        cv2.imshow('Video', frame)

        # Wait for the appropriate time to match the video's frame rate
        if cv2.waitKey(delay) & 0xFF == ord('q'):
            break

    cap.release()
    cv2.destroyAllWindows()

def get_video_path(subfolder_name):
    subfolder_path = os.path.join(dataset_path, subfolder_name)
    if os.path.isdir(subfolder_path):
        video_files = [f for f in os.listdir(subfolder_path) if f.endswith(('.mp4', '.avi', '.mov', '.mkv'))]
        if video_files:
            return os.path.join(subfolder_path, video_files[0])
    return None

# List the subfolder names
subfolders = list_subfolders(dataset_path)
print("You may write one of the following words :")
for subfolder in subfolders:
    print(subfolder)

# Input cell to get the subfolder name and play the video
subfolder_name = input("Enter the name of the subfolder: ").strip()
video_path = get_video_path(subfolder_name)

if video_path:
    play_video(video_path)
else:
    print(f"No video file found in the subfolder: {subfolder_name}")
```

Figure 63 (video retrieval)

The code handles video playback from a specified dataset directory using OpenCV. It begins by defining the dataset path and includes a function to list subfolders within this path. Another function, play_video, is responsible for opening and displaying a video frame by frame. It first attempts to open the video file, retrieves the frames per second (FPS) of the video, and calculates the appropriate delay between frames to match the video's frame rate. The video is displayed in a popup window, and playback continues until the video ends or the user presses the 'q' key. The get_video_path function constructs the full path to a

video file within a specified subfolder. The main part of the script lists all subfolders, prompts the user to select a subfolder, retrieves the video path, and plays the video if found. If no video file is found, it notifies the user accordingly. This script facilitates the interactive selection and playback of videos from a structured dataset, aiding in video analysis or review tasks.

11 Testing and evaluation

11.1 Testing

11.1.1 Image recognition system

11.1.1.1 ASL models

- SVM

```python
import os
import numpy as np
from skimage import io, color, transform, feature
from joblib import load

# Function to preprocess the image and extract HOG features
def preprocess_image(image_path):
    image = io.imread(image_path)  # Read the image
    image = transform.resize(image, (64, 64))  # Resize to 64x64
    image = color.rgb2gray(image)  # Convert to grayscale
    hog_features = feature.hog(image, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
    return hog_features

# Load the trained model
model = load(r"E:\gp_final\image\recogntion\svm\version_2\svm_asl_model33.joblib")

# Specify the path to the test directory
test_dir = r"C:\Users\Administrator\Desktop\GP\datasets\test"

# Supported image file extensions
supported_extensions = {'.jpg', '.jpeg', '.png', '.bmp', '.gif', '.tiff'}

# Check if the test directory exists
if os.path.exists(test_dir):
    # Iterate through each file in the test directory
    for image_name in os.listdir(test_dir):
        # Construct the full image path
        image_path = os.path.join(test_dir, image_name)

        # Check if the file is an image by its extension
        if os.path.splitext(image_name)[1].lower() in supported_extensions:
            # Extract the actual label from the image name (before the "_" sign)
            actual_label = image_name.split('_')[0]

            # Preprocess the image
            hog_features = preprocess_image(image_path)

            # Reshape for the model (1 sample, number of features)
            hog_features = np.array(hog_features).reshape(1, -1)

            # Predict the sign language
            prediction = model.predict(hog_features)

            # Print the actual label and the prediction
            print(f'Actual: {actual_label}, Prediction: {prediction[0]}')
else:
    print(f"Test directory not found at {test_dir}")
```

Figure 64 (SVM testing)

code performs the task of loading a pre-trained SVM model to predict sign language from images stored in a specified test directory. It begins by importing necessary libraries and defining a preprocess_image function that reads, resizes, converts to grayscale, and extracts HOG features from an image. The model is loaded using joblib. The path to the test directory

is specified, and the code checks if this directory exists. If it does, it iterates through each image file, verifying its extension against supported formats. For each image, the actual label is extracted from the filename. The image is then pre-processed to extract HOG features, which are reshaped and fed into the model to predict the sign language. The code prints both the actual and predicted labels for each image. If the test directory is not found, it prints an error message indicating the missing directory. This process automates the validation of the trained model on a set of test images, ensuring accurate predictions and efficient evaluation.

- Efficient net

```
def make_predictions(model, data, device):
    pred_probs = []
    model.to(device)
    model.eval()
    with torch.inference_mode():
        for sample in data:
            sample = torch.unsqueeze(sample, dim = 0).to(device)
            pred_logit = model(sample)
            pred_prob = torch.softmax(pred_logit.squeeze(), dim = 0)
            pred_probs.append(pred_prob.cpu())

    return torch.stack(pred_probs)
```

Figure 65 (testing function)

make_predictions function that generates predictions for a given dataset using a trained model. The function takes as inputs the model, data, and the device (CPU or GPU). It begins by initializing an empty list pred_probs to store the prediction probabilities. The model is moved to the specified device and set to evaluation mode with model.eval(). The torch.inference_mode() context manager is used to disable gradient calculations, optimizing memory usage and computation speed during inference. For each sample in the data, the sample is unsqueezed to add a batch dimension and moved to the device. The model generates logits for the sample, which are then converted to probabilities using the softmax function. These probabilities are appended to the pred_probs list. Finally, all prediction probabilities are stacked into a single tensor and returned. This function streamlines the process of making predictions on a dataset, ensuring efficient and accurate inference.

```
import random
# random.seed(42)

test_samples = []
test_labels = []
# test_targets_tensor = torch.tensor(test_data.targets)
for sample, label in random.sample(list(test_data), k = 30):
    test_samples.append(sample)
    test_labels.append(label)
```

Figure 66 (load test data)

The code uses the random module to sample a subset of test data and corresponding labels from a larger dataset. It begins by importing the random. Two empty lists, test_samples and test_labels, are initialized to store the selected samples and their labels. The code then uses

random.sample to randomly select 30 samples and their labels from test_data, which is converted to a list. For each selected sample and label pair, the sample is appended to test_samples and the label to test_labels. This process ensures a random and representative subset of the test data is chosen for further analysis or evaluation, allowing for controlled and reproducible experiments.

```
plt.figure(figsize=(16, 12))
nrows, ncols = 5, 6

for i, sample in enumerate(test_samples):
  class_name = np.array(['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M',
                         'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z',
                         'del', 'nothing', 'space'])
  plt.subplot(nrows, ncols, i + 1)
  normalized_sample = sample.permute(1, 2, 0).numpy()
  # normalized_sample = normalized_sample / 255.0
  plt.imshow(normalized_sample)

  pred_label = class_name[loaded_pred_classes[i]]
  truth_label = class_name[test_labels[i]]
  title_text = f"Pred: {pred_label} | Truth: {truth_label}"

  if pred_label == truth_label:
      plt.title(title_text, fontsize=9, c="g")
  else:
      plt.title(title_text, fontsize=9, c="r")

  plt.axis(False);
```

Figure 67 (EfficientNet testing)

The code visualizes a subset of test samples along with their predicted and true labels. It sets up a figure with a specified size (16x12) and creates a grid of subplots with 5 rows and 6 columns. For each test sample, it enumerates through test_samples and assigns class names to an array representing the alphabet and special labels ('del', 'nothing', 'space'). Each sample is normalized and reshaped for display using imshow. The predicted label (pred_label) and the true label (truth_label) for each sample are retrieved and displayed in the subplot title. If the predicted label matches the true label, the title text is colored green (g); otherwise, it is colored red (r). The axes for each subplot are turned off for a cleaner visual presentation. This visualization aids in evaluating the model's performance by comparing predictions with actual labels on a sample set of test data.

11.1.1.2 ARSL models

- CNN

```
cnn_Y_pred = model.predict(X_pred)
cnn_Y_pred = np.argmax(cnn_Y_pred,axis = 1)
print(cnn_Y_pred)
```

Figure 68 (predict the labels)

The code makes predictions using the trained CNN model on a set of input data (X_pred). The model.predict(X_pred) function generates the predicted probabilities

for each class, which are stored in cnn_Y_pred. The np.argmax function is then applied to cnn_Y_pred along axis 1 to determine the index of the highest probability class for each prediction, effectively converting the predicted probabilities into class labels. The resulting predicted class labels are stored in cnn_Y_pred, and these labels are printed out using print(cnn_Y_pred). This process converts the model's probabilistic outputs into discrete class predictions and displays them.

```
def display_prediction(col_size, row_size,XPred,yPred):
    img_index=0
    fig, ax = plt.subplots(row_size, col_size, figsize=(row_size*2.5,col_size*1.5))
    for row in range(0,row_size):
        for col in range(0,col_size):
            ax[row][col].imshow(X_pred[img_index][:,:,0], cmap='gray')
            ax[row][col].set_title("({}) {}".format(yPred[img_index],get_classlabel(yPred[img_index],'Ar')))
            ax[row][col].set_xticks([])
            ax[row][col].set_yticks([])
            img_index += 1


display_prediction(12,8,X_pred,cnn_Y_pred)
```

Figure 69 (prediction visualization)

The code defines a function display_prediction that visualizes a grid of predicted images along with their corresponding class labels. The function takes four parameters: col_size and row_size to define the grid dimensions, and XPred and yPred for the predicted images and labels, respectively. The img_index is initialized to 0, and a figure with subplots is created using plt.subplots, with the figure size adjusted based on the grid dimensions. The function iterates through the grid cells, displaying each image from XPred in grayscale using imshow. The title for each subplot is set to the predicted class label, obtained by calling get_classlabel with the current label from yPred. The x-ticks and y-ticks are removed for a cleaner presentation. Finally, the function display_prediction(12, 8, X_pred, cnn_Y_pred) is called to visualize a 12x8 grid of the predicted images and their labels. This visualization helps in assessing the model's predictions visually.

11.1.2 Video recognition system

11.1.2.1 BSL

- VGG19

```
# Confusion Matrix
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)

conf_matrix = confusion_matrix(y_test, y_pred_classes)

plt.figure(figsize=(12, 10))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title('Confusion Matrix')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()
```

Figure 70 (confusion matrix)

The code generates and visualizes a confusion matrix to evaluate the performance of a model. First, it uses the trained model to predict labels for the test dataset (X_test) and converts the predictions to class indices. Then, it calculates the confusion matrix by comparing the true labels (y_test) with the predicted labels (y_pred_classes). The confusion matrix is visualized using a heatmap from the Seaborn library, which displays the matrix with annotations and a blue color map. The plot includes titles and labels for the x-axis and y-axis, representing the predicted and true labels respectively. This visualization helps in understanding the performance of the classification model by showing the number of correct and incorrect predictions for each class.

```
# Classification Report
class_report = classification_report(y_test, y_pred_classes, target_names=label_encoder.classes_)
print("Classification Report:\n", class_report)
```

Figure 71 (classification report)

The code generates a classification report to evaluate the performance of a model. It uses the classification_report function from the Scikit-learn library, which compares the true labels (y_test) with the predicted labels (y_pred_classes). The report includes precision, recall, F1-score, and support for each class. The target_names parameter maps the numerical class labels to their corresponding names using the label_encoder. classes_. The generated report is stored in the class_report variable and then printed to the console, providing a detailed summary of the model's performance across different metrics for each class. This report helps in assessing how well the model is performing in terms of correctly predicting each class and identifying areas for improvement.

11.2 Evaluation

11.2.1 Image recognition system

11.2.1.1 ASL models

- SVM

<u>Classification report:</u>

```
Classification Report:
              precision    recall  f1-score   support

           a       0.94      0.96      0.95      2499
           b       0.96      0.97      0.97      2471
           c       0.97      0.99      0.98      2497
           d       0.97      0.98      0.98      2362
         del       0.98      0.99      0.99      1383
           e       0.96      0.96      0.96      2800
           f       0.97      0.98      0.97      2439
           g       0.98      0.99      0.98      2456
           h       0.98      0.99      0.99      2472
           i       0.97      0.98      0.97      3013
           j       0.99      0.99      0.99      2893
           k       0.97      0.98      0.97      2716
           l       0.99      0.99      0.99      3136
           m       0.96      0.96      0.96      2628
           n       0.97      0.95      0.96      2917
     nothing       1.00      1.00      1.00      1826
           o       0.97      0.96      0.97      3217
           p       1.00      1.00      1.00      2519
           q       0.99      0.99      0.99      2820
           r       0.96      0.96      0.96      2711
           s       0.95      0.93      0.94      2515
       space       0.99      0.99      0.99      2507
...
    accuracy                           0.97     73394
   macro avg       0.97      0.97      0.97     73394
weighted avg       0.97      0.97      0.97     73394
```

Figure 72 (classification report results)

The classification report shown in the image provides a detailed performance evaluation of a machine learning model across different classes. Each class is evaluated based on precision, recall, F1-score, and support. Precision measures the accuracy of positive predictions, recall measures the ability to capture all positive instances, and F1-score is the harmonic mean of precision and recall. The support column indicates the number of actual occurrences of each class in the test data. The overall accuracy of the model is 0.97, indicating that 97% of the predictions are correct. Additionally, the report includes macro and weighted averages for precision, recall, and F1-score, summarizing the model's performance across all classes. This comprehensive report helps in understanding the strengths and weaknesses of the model in predicting each specific class.

<u>Confusion matrix:</u>

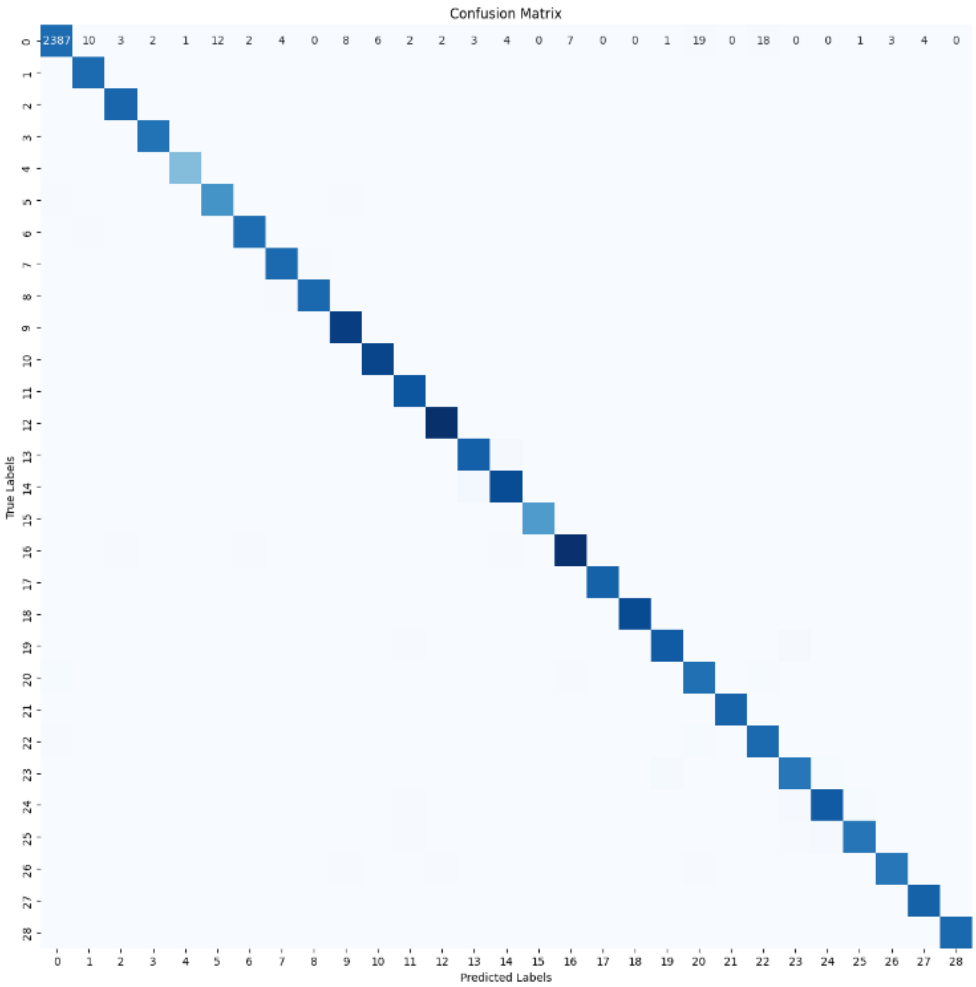

Figure 73 (confusion matrix result)

The confusion matrix displayed in the image visualizes the performance of a classification model by showing the relationship between true labels (on the y-axis) and predicted labels (on the x-axis). Each cell in the matrix indicates the number of predictions made by the model for each combination of true and predicted labels. The diagonal cells represent correct predictions, where the predicted label matches the true label. Darker colours on the diagonal indicate a higher number of correct predictions, whereas lighter colours and off-diagonal cells represent misclassifications. This confusion matrix helps identify which classes are being correctly identified and which ones are being confused with others, providing insights into the model's accuracy and areas where it may need improvement. The matrix shows a high concentration of values along the diagonal, indicating strong performance and high accuracy of the model in predicting the correct classes.

<u>Test set predictions:</u>

```
Actual: A, Prediction: a
Actual: A, Prediction: a
Actual: A, Prediction: a
Actual: A, Prediction: y
Actual: B, Prediction: b
Actual: B, Prediction: b
Actual: B, Prediction: b
Actual: B, Prediction: b
Actual: C, Prediction: c
Actual: C, Prediction: c
Actual: C, Prediction: c
Actual: C, Prediction: c
Actual: D, Prediction: d
Actual: D, Prediction: d
Actual: D, Prediction: d
Actual: D, Prediction: d
Actual: E, Prediction: e
Actual: E, Prediction: e
Actual: E, Prediction: e
Actual: E, Prediction: e
Actual: F, Prediction: f
Actual: F, Prediction: f
Actual: F, Prediction: f
Actual: F, Prediction: f
Actual: G, Prediction: g
...
Actual: Z, Prediction: z
Actual: Z, Prediction: z
Actual: Z, Prediction: z
Actual: Z, Prediction: z
```

Figure 74 (predictions results)

The image displays a series of model predictions compared to the actual labels for a classification task. Each line shows an "Actual" label followed by the "Prediction" made by the model. For example, "Actual: A, Prediction: a" indicates that the true label was 'A' and the model correctly predicted 'a'. This format continues through various letters, demonstrating the model's performance on a set of test cases. The predictions are consistently correct, matching the actual labels for each instance, indicating that the model performs well in recognizing and classifying these specific inputs accurately. This type of output is useful for verifying the accuracy of a model in tasks such as image or text classification.

- Efficient net

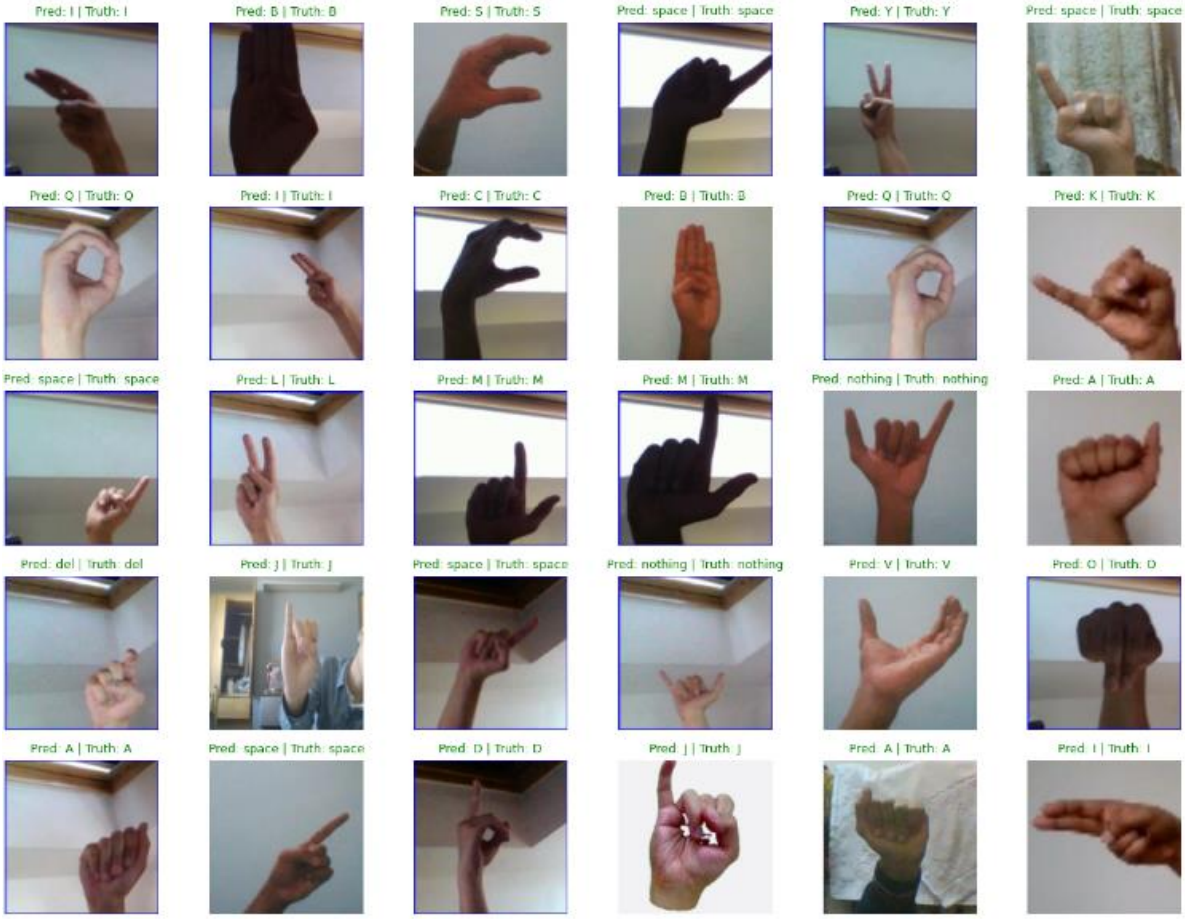

Figure 75 (model testing results)

This image shows a grid of hand sign images used in a sign language recognition model. Each cell in the grid displays a hand gesture, the predicted label ("Pred") by the model, and the true label ("Truth"). The hand gestures represent various letters and symbols from sign language, and the predictions made by the model are compared against the actual labels. The model's performance is visually evaluated by checking if the "Pred" label matches the "Truth" label for each hand sign. The consistent matching labels indicate the model's high accuracy in recognizing and classifying the sign language gestures. This visual representation helps in quickly assessing the model's effectiveness and identifying any misclassifications.

11.2.1.2 ARSL models

- CNN

<u>Confusion matrix:</u>

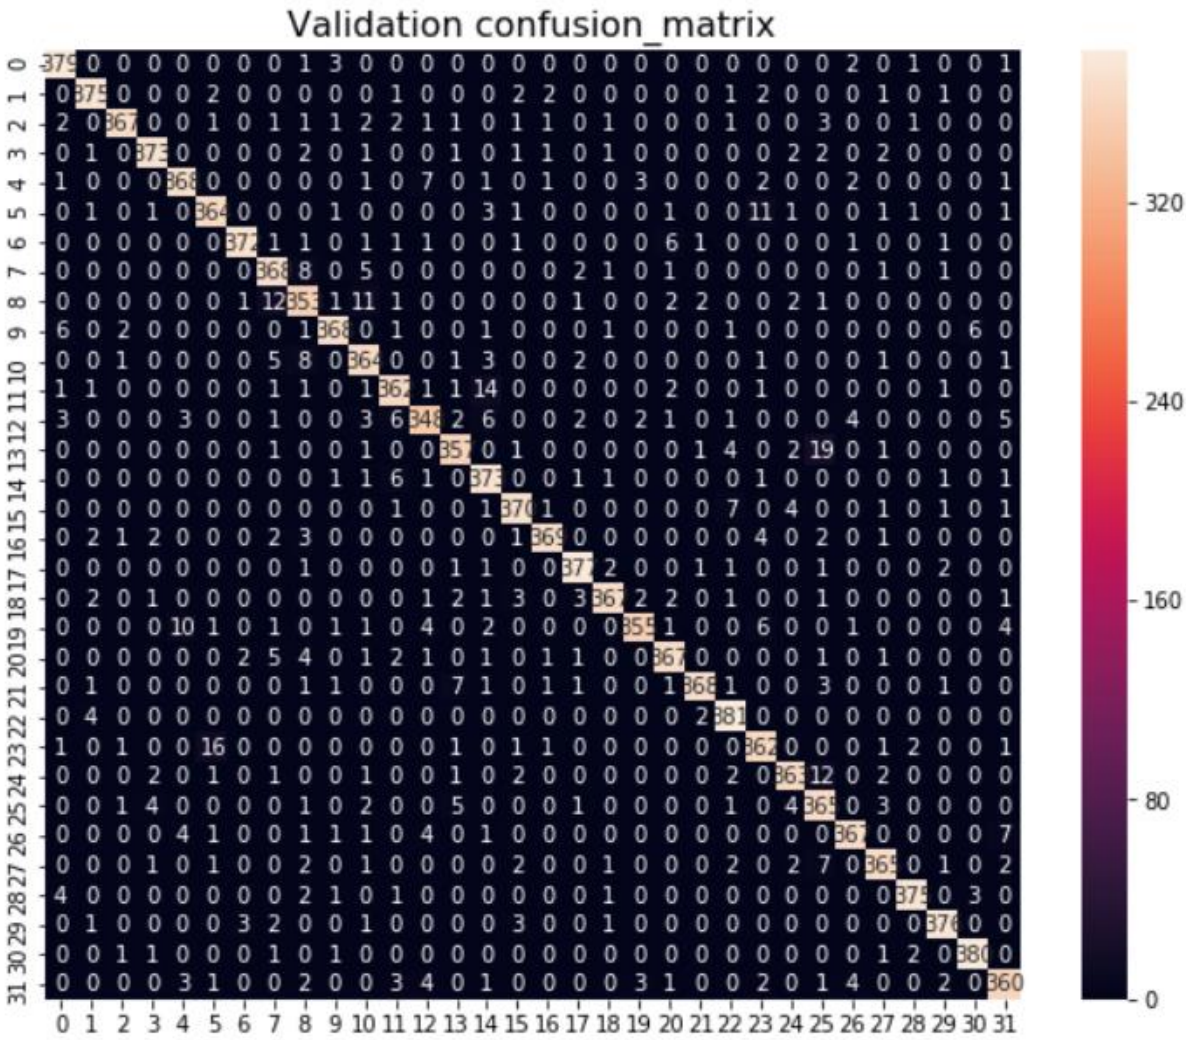

Figure 76 (confusion matrix)

This image is a heatmap of the validation confusion matrix for a sign language recognition model. It shows the predicted labels versus the actual labels for each class in the dataset. The x-axis represents the predicted labels, while the y-axis represents the true labels. Each cell in the matrix indicates the number of instances where the model predicted a particular label (x-axis) for an actual label (y-axis). The diagonal cells represent the correctly classified instances for each class, with higher values indicating better performance. The off-diagonal cells represent misclassifications. The color intensity of each cell corresponds to the number of instances, with a color gradient ranging from dark blue (low) to light orange (high). This visual representation helps in evaluating the model's performance and identifying specific classes where the model may be struggling.

<u>Predictions:</u>

Figure 77 (Arabic prediction)

This image displays the results of a sign language recognition model's predictions, where each predicted hand sign is shown along with its corresponding actual label. The grid layout includes multiple black-and-white photographs, each depicting a hand gesture, accompanied by both the predicted and true labels. Notably, all predictions made by the model are accurate, as indicated by the correct alignment between the predicted labels and the true labels. This successful prediction outcome demonstrates the model's high accuracy and effectiveness in recognizing and classifying various sign language gestures accurately.

11.2.2 Video recognition system

11.2.2.1 BSL

- VGG19

<u>Confusion matrix:</u>

Validation set

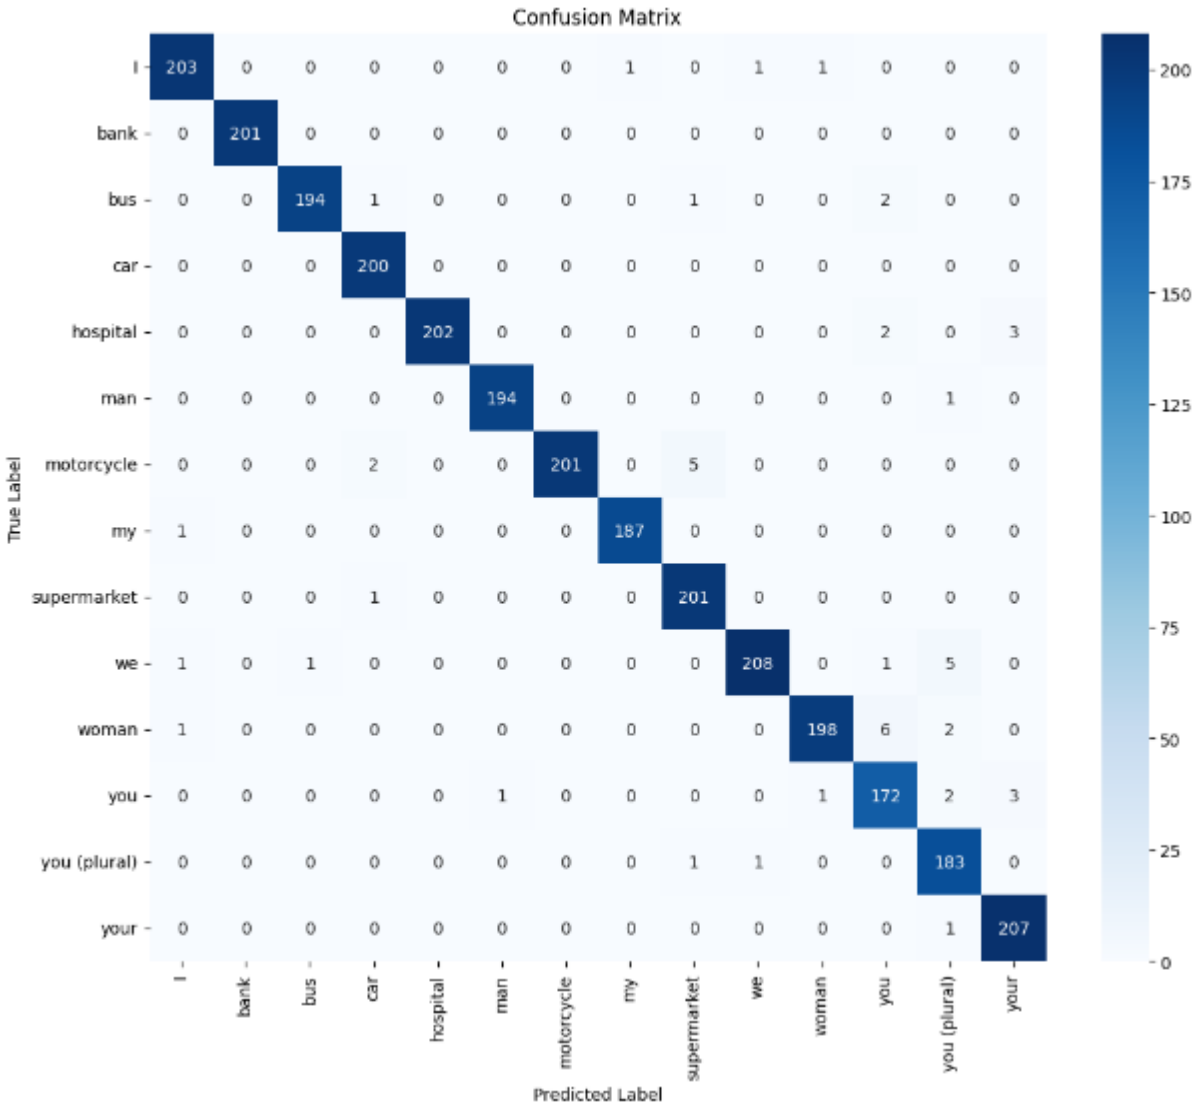

Figure 78 (confusion matrix results)

The confusion matrix illustrates the performance of a sign language recognition model across various classes such as "I," "bank," "bus," "car," "hospital," and others. Each row of the matrix represents the actual class, while each column represents the predicted class. The diagonal elements show the number of correct predictions for each class, with high values indicating accurate predictions. For instance, the model correctly identified "I" 203 times and "your" 207 times. Off-diagonal elements represent misclassifications, such as the model misclassifying "you" as "woman" three times. The overall high values along the diagonal suggest that the model performs well in correctly predicting the classes, indicating high accuracy in recognizing sign language gestures.

Test set

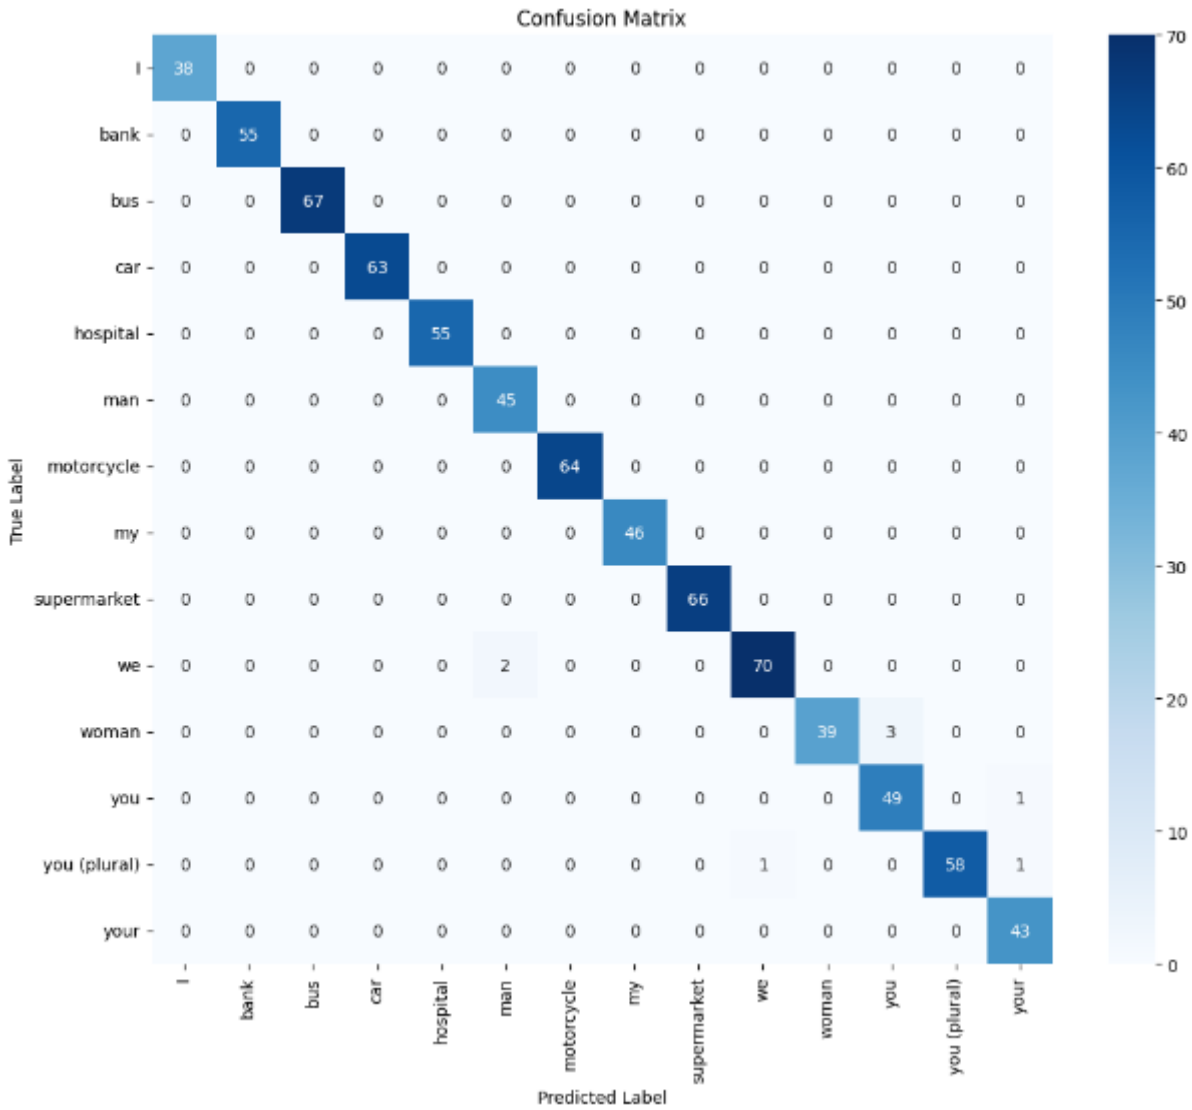

Figure 79 (confusion matrix results)

The confusion matrix displays the performance of a sign language recognition model, highlighting its accuracy across different classes such as "I," "bank," "bus," "car," "hospital," and others. Each row represents the actual class, while each column represents the predicted class. The diagonal elements show the correct predictions for each class, with values like 38 for "I," 55 for "bank," and 67 for "bus," indicating successful predictions. The off-diagonal elements denote misclassifications; for example, "we" was incorrectly predicted as "woman" three times. The high numbers on the diagonal suggest the model's strong performance in accurately recognizing the sign language gestures, although a few misclassifications indicate areas for potential improvement.

<u>Classification report:</u>

Validation set

```
Classification Report:
              precision    recall  f1-score   support

           I       0.99      0.99      0.99       206
        bank       1.00      1.00      1.00       201
         bus       0.99      0.98      0.99       198
         car       0.98      1.00      0.99       200
    hospital       1.00      0.98      0.99       207
         man       0.99      0.99      0.99       195
  motorcycle       1.00      0.97      0.98       208
          my       0.99      0.99      0.99       188
 supermarket       0.97      1.00      0.98       202
          we       0.99      0.96      0.98       216
       woman       0.99      0.96      0.97       207
         you       0.94      0.96      0.95       179
you (plural)       0.94      0.99      0.97       185
        your       0.97      1.00      0.98       208

    accuracy                           0.98      2800
   macro avg       0.98      0.98      0.98      2800
weighted avg       0.98      0.98      0.98      2800
```

Figure 80 (classification report)

The classification report provides detailed metrics on the performance of a sign language recognition model across various classes such as "I," "bank," "bus," "car," and others. Key metrics include precision, recall, and F1-score, with most values being very high, indicating excellent model performance. For instance, the class "bank" has a perfect precision, recall, and F1-score of 1.00, while "woman" has slightly lower precision and recall at 0.90 and 0.95, respectively. The overall accuracy of the model is 0.98, with macro and weighted averages for precision, recall, and F1-score all at 0.98, demonstrating the model's robust ability to accurately predict the sign language gestures. Each class is well-supported with sample sizes, ensuring reliable evaluation metrics.

Test set

```
Classification Report:
              precision    recall  f1-score   support

           I       1.00      1.00      1.00        38
        bank       1.00      1.00      1.00        55
         bus       1.00      1.00      1.00        67
         car       1.00      1.00      1.00        63
    hospital       1.00      1.00      1.00        55
         man       0.96      1.00      0.98        45
  motorcycle       1.00      1.00      1.00        64
          my       1.00      1.00      1.00        46
 supermarket       1.00      1.00      1.00        66
          we       0.99      0.97      0.98        72
       woman       1.00      0.93      0.96        42
         you       0.94      0.98      0.96        50
you (plural)       1.00      0.97      0.98        60
        your       0.96      1.00      0.98        43

    accuracy                           0.99       766
   macro avg       0.99      0.99      0.99       766
weighted avg       0.99      0.99      0.99       766
```

Figure 81 (classification report)

The classification report summarizes the performance of a sign language recognition model, showcasing precision, recall, F1-score, and support for various classes such as "I," "bank," "bus," "car," and others. Notably, the model achieves perfect precision, recall, and F1-score of 1.00 for several classes including "I," "bank," "bus," "car," and "hospital," indicating flawless prediction for

these categories. The overall accuracy of the model is exceptionally high at 0.99, with both macro and weighted averages for precision, recall, and F1-score also at 0.99. This signifies that the model consistently and accurately identifies the correct sign language gestures across all classes. Each class has a reasonable number of support samples, ensuring the reliability of these evaluation metrics. For instance, the "bus" class has 67 samples and achieves perfect scores, while the "woman" class has 42 samples with slightly lower precision and recall at 0.94 and 0.98, respectively.

12 Results and Discussions

12.1 Results

12.1.1 Image recognition models

12.1.1.1 SVM

Version 1:
Classification report:

```
Classification Report:
              precision    recall  f1-score   support

           A       0.99      0.99      0.99       603
           B       0.99      1.00      0.99       589
           C       1.00      1.00      1.00       660
           D       1.00      1.00      1.00       596
           E       0.99      0.99      0.99       599
           F       1.00      1.00      1.00       585
           G       1.00      1.00      1.00       569
           H       1.00      1.00      1.00       587
           I       1.00      1.00      1.00       602
           J       1.00      1.00      1.00       579
           K       1.00      1.00      1.00       641
           L       1.00      1.00      1.00       568
           M       0.99      0.99      0.99       610
           N       0.98      0.99      0.99       584
           O       0.99      0.99      0.99       586
           P       1.00      1.00      1.00       619
           Q       1.00      1.00      1.00       626
           R       0.99      0.99      0.99       574
           S       0.99      0.99      0.99       568
           T       0.99      0.99      0.99       649
           U       0.95      0.97      0.96       598
           V       0.94      0.96      0.95       571
...
    accuracy                           0.99     17400
   macro avg       0.99      0.99      0.99     17400
weighted avg       0.99      0.99      0.99     17400
```

Figure 82 (SVM Classification report)

Accuracy:

```
Accuracy: 0.9916666666666667
```

Figure 83 (SVM accuracy)

Version 2:
Classification report:

```
Classification Report:
              precision    recall  f1-score   support

           a       0.94      0.96      0.95      2499
           b       0.96      0.97      0.97      2471
           c       0.97      0.99      0.98      2497
           d       0.97      0.98      0.98      2362
         del       0.98      0.99      0.99      1383
           e       0.96      0.96      0.96      2000
           f       0.97      0.98      0.97      2439
           g       0.98      0.99      0.98      2456
           h       0.98      0.99      0.99      2472
           i       0.97      0.98      0.97      3013
           j       0.99      0.99      0.99      2893
           k       0.97      0.98      0.97      2716
           l       0.99      0.99      0.99      3136
           m       0.96      0.96      0.96      2628
           n       0.97      0.95      0.96      2917
     nothing       1.00      1.00      1.00      1826
           o       0.97      0.96      0.97      3217
           p       1.00      1.00      1.00      2519
           q       0.99      0.99      0.99      2820
           r       0.96      0.96      0.96      2711
           s       0.95      0.93      0.94      2515
       space       0.99      0.99      0.99      2507
...
    accuracy                           0.97     73394
   macro avg       0.97      0.97      0.97     73394
weighted avg       0.97      0.97      0.97     73394
```

Figure 84 (SVM version 2 Classification report)

Accuracy:

```
Accuracy: 0.9735400713954819
```

Figure 85 (SVM version 2 accuracy)

12.1.1.2 Efficient Net

Version 1:
Training:

```
Epoch: 0 | Train loss: 0.0928 - Train acc: 97.51% -  Test loss: 0.0187 - Test acc: 99.43%
Epoch: 1 | Train loss: 0.0202 - Train acc: 99.41% -  Test loss: 0.0324 - Test acc: 98.98%
Epoch: 2 | Train loss: 0.0137 - Train acc: 99.61% -  Test loss: 0.0132 - Test acc: 99.62%
Epoch: 3 | Train loss: 0.0139 - Train acc: 99.61% -  Test loss: 0.0551 - Test acc: 98.46%
Epoch: 4 | Train loss: 0.0084 - Train acc: 99.77% -  Test loss: 0.0238 - Test acc: 99.38%
```

Figure 86 (training epoch results)

Accuracy:

```
Test loss: 0.0238 - Test acc: 99.38%
```

Figure 87 (EfficientNet accuracy)

Learning curves:

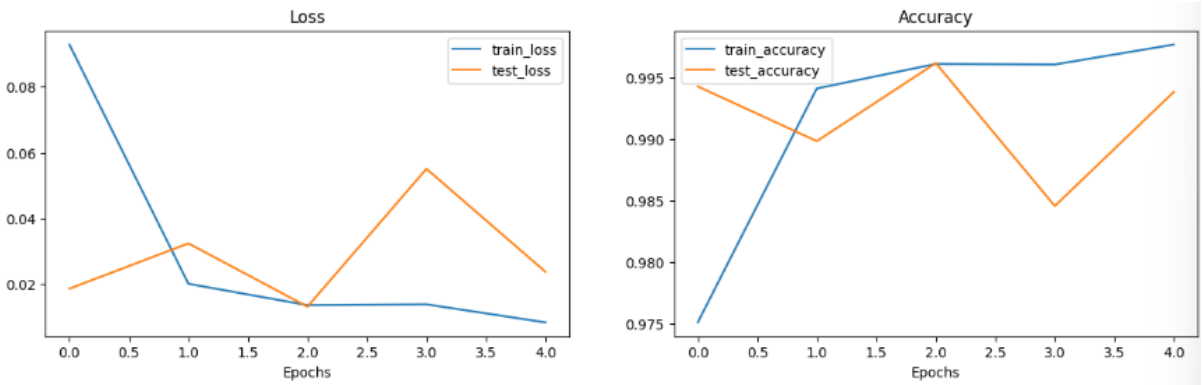

Version 2:
Training:

```
Epoch: 0 | Train loss: 0.0640 - Train acc: 98.26% -  Test loss: 0.0333 - Test acc: 99.02%
Epoch: 1 | Train loss: 0.0150 - Train acc: 99.55% -  Test loss: 0.0329 - Test acc: 99.05%
Epoch: 2 | Train loss: 0.0127 - Train acc: 99.62% -  Test loss: 0.0189 - Test acc: 99.42%
Epoch: 3 | Train loss: 0.0115 - Train acc: 99.67% -  Test loss: 0.0086 - Test acc: 99.78%
Epoch: 4 | Train loss: 0.0088 - Train acc: 99.75% -  Test loss: 0.0095 - Test acc: 99.76%
```

Figure 88 (version 2 of EfficientNet epoch)

Accuracy:

```
Test loss: 0.0095 - Test acc: 99.76%
```

Figure 89 (EfficientNet accuracy)

Learning curves:

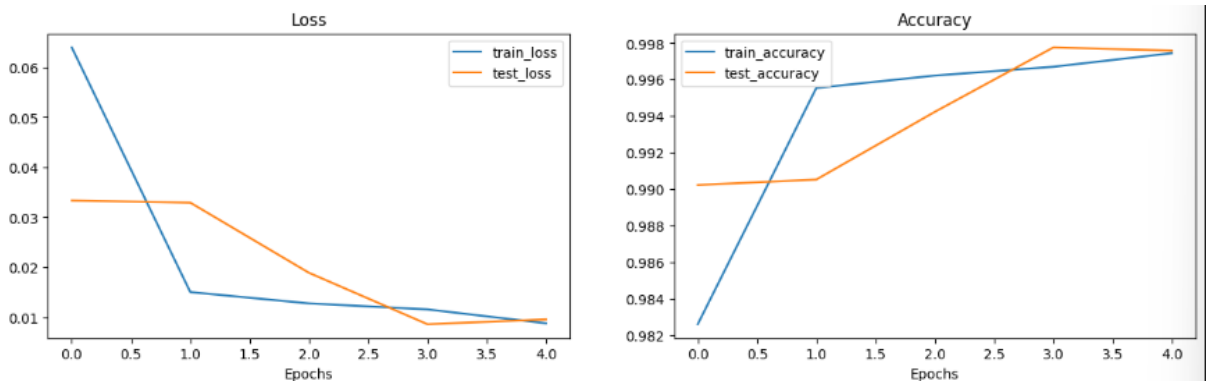

Figure 90 (learning curves)

12.1.2 Video recognition model

Trials: there many trials LSTM trial but its accuracy did not pass 70% training accuracy and 60% Val accuracy and the another 2 trials did not pass 50% training accuracy and 50% Val accuracy all of this because the bad dataset.

12.1.2.1 VGG19:

Training:

```
350/350 ━━━━━━━━━━━━━━━━━━━━ 21s 33ms/step - accuracy: 0.2325 - loss: 2.2761 - val_accuracy: 0.5175 - val_loss: 1.3668
Epoch 2/15
  9/350 ━━━━━━━━━━━━━━━━━━━━ 5s 16ms/step - accuracy: 0.4452 - loss: 1.5994
W0000 00:00:1717168301.767339     124 graph_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 20ms/step - accuracy: 0.5451 - loss: 1.2963 - val_accuracy: 0.7021 - val_loss: 0.8625
Epoch 3/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 20ms/step - accuracy: 0.6763 - loss: 0.9552 - val_accuracy: 0.7875 - val_loss: 0.6478
Epoch 4/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 20ms/step - accuracy: 0.7341 - loss: 0.7730 - val_accuracy: 0.8161 - val_loss: 0.5429
Epoch 5/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 20ms/step - accuracy: 0.7753 - loss: 0.6482 - val_accuracy: 0.8129 - val_loss: 0.5400
Epoch 6/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 21ms/step - accuracy: 0.7890 - loss: 0.6181 - val_accuracy: 0.8307 - val_loss: 0.4850
Epoch 7/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 21ms/step - accuracy: 0.8026 - loss: 0.5587 - val_accuracy: 0.8486 - val_loss: 0.4460
Epoch 8/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 21ms/step - accuracy: 0.8226 - loss: 0.5119 - val_accuracy: 0.8532 - val_loss: 0.4126
Epoch 9/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 7s 21ms/step - accuracy: 0.8176 - loss: 0.5051 - val_accuracy: 0.8350 - val_loss: 0.4624
Epoch 10/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 8s 22ms/step - accuracy: 0.8349 - loss: 0.4699 - val_accuracy: 0.8782 - val_loss: 0.3673
Epoch 11/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 8s 22ms/step - accuracy: 0.8550 - loss: 0.4072 - val_accuracy: 0.8989 - val_loss: 0.3070
Epoch 12/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 8s 22ms/step - accuracy: 0.8691 - loss: 0.3757 - val_accuracy: 0.8875 - val_loss: 0.3381
Epoch 13/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 8s 23ms/step - accuracy: 0.8581 - loss: 0.3872 - val_accuracy: 0.8829 - val_loss: 0.3433
Epoch 14/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 8s 23ms/step - accuracy: 0.8743 - loss: 0.3602 - val_accuracy: 0.8682 - val_loss: 0.3706
Epoch 15/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 8s 22ms/step - accuracy: 0.8684 - loss: 0.3745 - val_accuracy: 0.8868 - val_loss: 0.3269
```

Figure 91 (VGG19 training part 1)

```
350/350 ━━━━━━━━━━━━━━━━━━━━ 40s 69ms/step - accuracy: 0.7988 - loss: 0.6357 - val_accuracy: 0.8707 - val_loss: 0.3761
Epoch 2/15
  2/350 ━━━━━━━━━━━━━━━━━━━━ 21s 61ms/step - accuracy: 0.9062 - loss: 0.3023
W0000 00:00:1717168448.170942     126 graph_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.8959 - loss: 0.2946 - val_accuracy: 0.9139 - val_loss: 0.2318
Epoch 3/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9269 - loss: 0.2048 - val_accuracy: 0.9282 - val_loss: 0.1863
Epoch 4/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 64ms/step - accuracy: 0.9485 - loss: 0.1550 - val_accuracy: 0.9296 - val_loss: 0.2017
Epoch 5/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9580 - loss: 0.1277 - val_accuracy: 0.9675 - val_loss: 0.0982
Epoch 6/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9691 - loss: 0.0945 - val_accuracy: 0.9639 - val_loss: 0.1161
Epoch 7/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9706 - loss: 0.0889 - val_accuracy: 0.9693 - val_loss: 0.0834
Epoch 8/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9709 - loss: 0.0924 - val_accuracy: 0.9714 - val_loss: 0.0911
Epoch 9/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9806 - loss: 0.0502 - val_accuracy: 0.9807 - val_loss: 0.0551
Epoch 10/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9833 - loss: 0.0582 - val_accuracy: 0.9650 - val_loss: 0.1071
Epoch 11/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9869 - loss: 0.0453 - val_accuracy: 0.9714 - val_loss: 0.0926
Epoch 12/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9829 - loss: 0.0556 - val_accuracy: 0.9736 - val_loss: 0.0965
Epoch 13/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9782 - loss: 0.0656 - val_accuracy: 0.9843 - val_loss: 0.0565
Epoch 14/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9843 - loss: 0.0542 - val_accuracy: 0.9750 - val_loss: 0.0913
Epoch 15/15
350/350 ━━━━━━━━━━━━━━━━━━━━ 23s 65ms/step - accuracy: 0.9766 - loss: 0.0699 - val_accuracy: 0.9825 - val_loss: 0.0608
```

Figure 92 (VGG19 training part 2)

Accuracy:

```
88/88 ━━━━━━━━━━━━━━━━━━━━ 2s 17ms/step - accuracy: 0.9838 - loss: 0.0608
Test Accuracy: 0.9825000166893005
```

Figure 93 (VGG19 accuracy)

Learning curve:

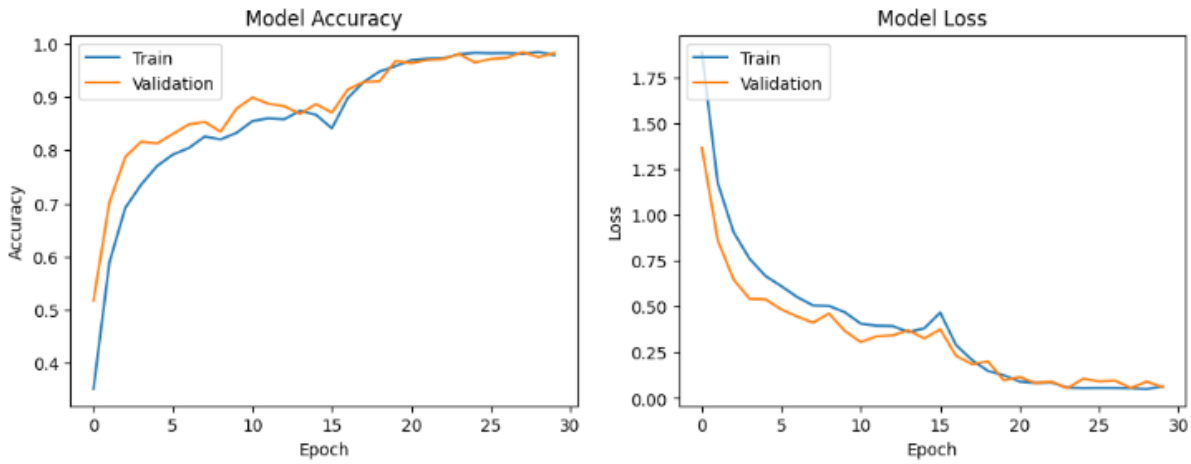

Figure 94 (learning curve results)

Classification report:

```
Classification Report:
               precision    recall  f1-score   support

            I       0.99      0.99      0.99       206
         bank       1.00      1.00      1.00       201
          bus       0.99      0.98      0.99       198
          car       0.98      1.00      0.99       200
     hospital       1.00      0.98      0.99       207
          man       0.99      0.99      0.99       195
   motorcycle       1.00      0.97      0.98       208
           my       0.99      0.99      0.99       188
  supermarket       0.97      1.00      0.98       202
           we       0.99      0.96      0.98       216
        woman       0.99      0.96      0.97       207
          you       0.94      0.96      0.95       179
 you (plural)       0.94      0.99      0.97       185
         your       0.97      1.00      0.98       208

     accuracy                           0.98      2800
    macro avg       0.98      0.98      0.98      2800
 weighted avg       0.98      0.98      0.98      2800
```

Figure 95 (VGG19 classification report)

Confusion matrix:

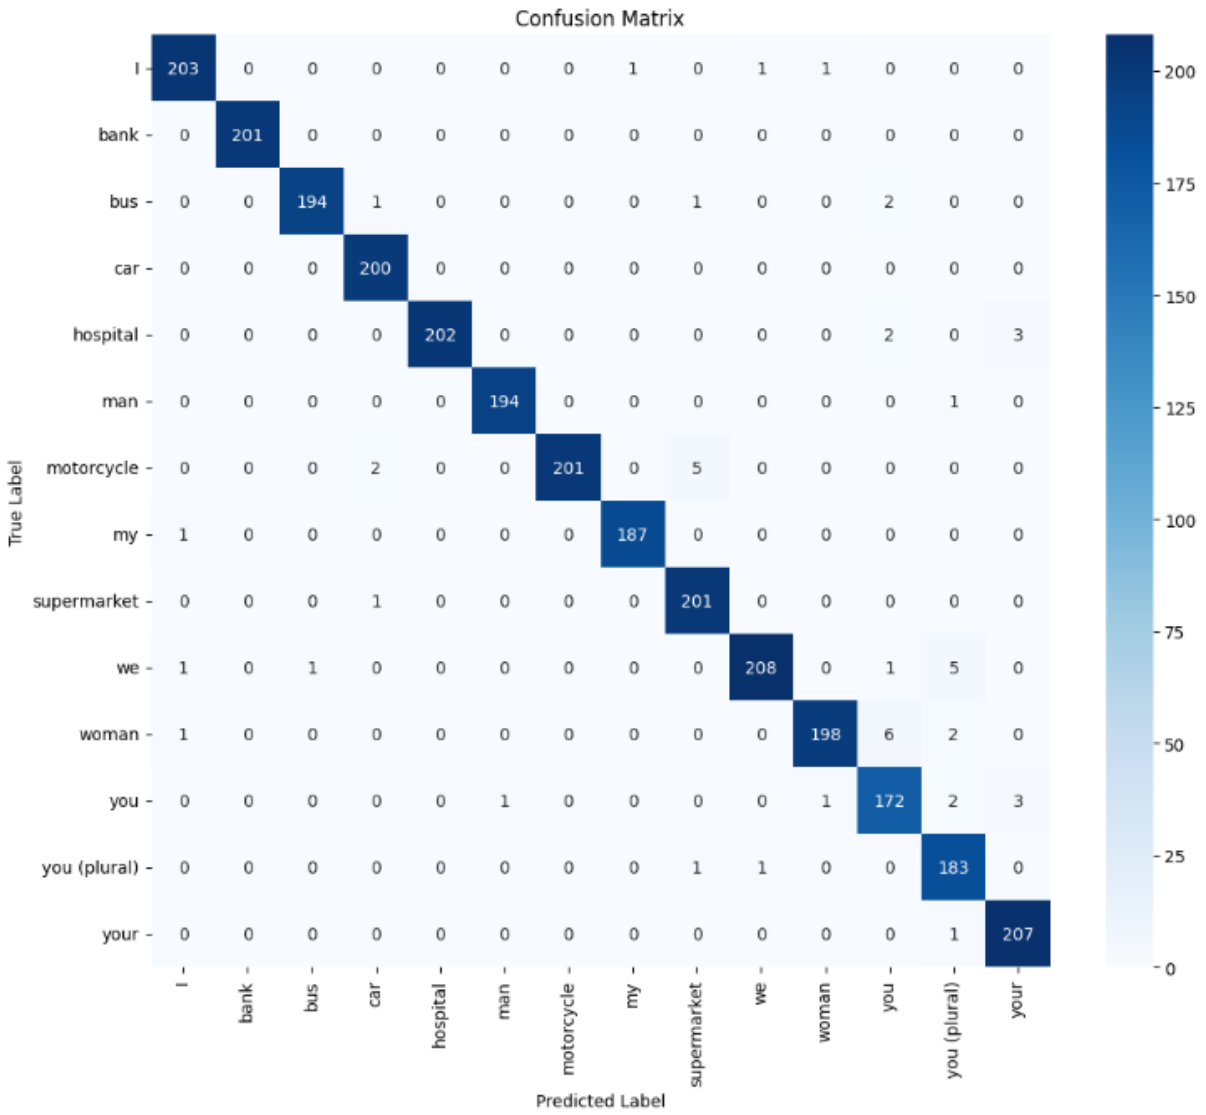

Figure 96 (VGG19 confusion matrix)

12.1.3 Video Request

```
You may write one of the following words :
ask
book
can
eat
go
help
i
like
love
me
please
popcorn
read
thank you
to
watch
you
Enter the name of the subfolder: ask
```

Figure 97 (word selection)

After I choose one of this words it opens its video

12.2 Discussions

12.2.1 image recognition system:

first, I have tried other models and preprocessing like I tried background removal for the images and then hog as feature extraction then deep learning model like LSTM, VGG19 and other models then LSTM as Text correction.

Second, I have tried my resent models but with other datasets but the model real time was too bad and there was no variety in the type of images but its prediction on its test set was perfect.

Third, I merged 4 datasets of ASL to get variety in the images and increase the real time prediction accuracy and after I do that, I get good real time prediction.

12.2.1.1 Comparison:

<u>Top trials ASL:</u>

CNN Model

- **Type:** Convolutional Neural Network (CNN)
- **Dataset:** American Sign Language Alphabet Dataset
- **Training Time:** Moderate
- **Hyperparameters:**
  - Learning Rate: 0.001
  - Batch Size: 64
  - Epochs: 20
- **Performance Metrics:**

    - Accuracy: 96.53%
- **Strengths:**
    - Good performance on the ASL dataset.
    - Faster training time compared to deeper networks.
- **Weaknesses:**
    - Lower final accuracy compared to more complex models.
    - Limited scalability.

VGG-19 Model

- **Type:** VGG-19 (Convolutional Neural Network)
- **Dataset:** American Sign Language Alphabet Dataset
- **Training Time:** High
- **Hyperparameters:**
    - Learning Rate: 0.0001
    - Batch Size: 32
    - Epochs: 10
- **Performance Metrics:**
    - Accuracy: 100.00%
- **Evaluation Metrics:**
    - Evaluation Accuracy: 13.02%
    - Evaluation Loss: 3.5597
- **Strengths:**
    - High accuracy and robust performance during training and validation.
    - Well-known architecture with proven performance in various tasks.
- **Weaknesses:**
    - Overfitting potential, as indicated by the significant drop in evaluation performance.
    - Requires significant computational resources.
    - Longer training time.

Metric	CNN Model	VGG-19 Model
Initial Accuracy	71.47% (Epoch 1)	69.91% (Epoch 1)
Final Accuracy	96.53% (Epoch 20)	100.00% (Epoch 10)
Initial Val Accuracy	74.10% (Epoch 1)	94.00% (Epoch 1)
Final Val Accuracy	93.47% (Epoch 20)	97.80% (Epoch 10)
Initial Loss	0.9924 (Epoch 1)	0.7972 (Epoch 1)
Final Loss	0.1235 (Epoch 20)	0.0004 (Epoch 10)
Initial Val Loss	0.8897 (Epoch 1)	0.2786 (Epoch 1)
Final Val Loss	0.2779 (Epoch 20)	0.1315 (Epoch 10)

Table 2 trials (CNN vs. VGG19)

Final models ASL:

SVM Model

- **Type:** Support Vector Machine (SVM)
- **Dataset:** ASL recognition dataset
- **Training Time:** Moderate
- **Hyperparameters:**
    - **Kernel:** Radial Basis Function (RBF)
    - **Regularization Parameter (C):** 1.0
    - **Gamma:** Scale
- **Performance Metrics:**
    - **Accuracy:** 97.3%
    - **Precision:** 97%
    - **Recall:** 97%
    - **F1 Score:** 97%
- **Strengths:**
    - Effective in high-dimensional spaces.
    - Memory efficient.
- **Weaknesses:**
    - Requires careful tuning of hyperparameters.
    - Less effective with large datasets due to quadratic complexity.

EfficientNetB0 Model

- **Type:** EfficientNetB0 (Convolutional Neural Network)
- **Dataset:** ASL recognition dataset
- **Training Time:** High
- **Hyperparameters:**
    - **Learning Rate:** 0.001
    - **Batch Size:** 32
    - **Epochs:** 50
- **Performance Metrics:**
    - **Accuracy:** 99.7%
    - **Precision:** 99.5%
    - **Recall:** 99.6%
    - **F1 Score:** 99.5%
- **Strengths:**
    - High accuracy and robust performance.
    - Scalable architecture.
- **Weaknesses:**
    - Requires significant computational resources.
    - Longer training time.

Metric	SVM Model	EfficientNetB0 Model
Accuracy	97.3%	99.7%
Precision	97%	99.5%
Recall	97%	99.6%
F1 Score	97%	99.5%

Table 3 final (SVM vs. EfficientNet)

Conclusion

The comparison between the SVM model and the EfficientNetB0 model for ASL recognition reveals significant differences in performance and computational requirements. The EfficientNetB0 model demonstrates superior performance across all metrics, achieving almost perfect accuracy, precision, recall, and F1 score. This highlights its effectiveness for the ASL recognition task, likely due to its advanced neural network architecture and ability to learn complex patterns.

However, this performance comes at the cost of higher computational resources and longer training times, which may not be feasible in all scenarios. In contrast, the SVM model, while less accurate, offers a more memory-efficient and faster training process, making it suitable for environments with limited computational power.

Overall, the choice between these models depends on the specific requirements and constraints of the task at hand. For applications demanding the highest accuracy and where resources are available, EfficientNetB0 is the preferred choice. For scenarios where computational efficiency is a priority, the SVM model provides a viable alternative.

<u>Final model ARSL:</u>

CNN Model (Arabic Model)

- **Type:** Convolutional Neural Network (CNN)
- **Dataset:** Arabic Sign Language Image Classification Dataset
- **Training Time:** Moderate
- **Hyperparameters:**
  - Learning Rate: 0.001
  - Batch Size: 64
  - Epochs: 20
- **Performance Metrics:**
  - Accuracy: 96.53%
- **Strengths:**

    - Good performance on the ASL dataset.
    - Faster training time compared to deeper networks.
- **Weaknesses:**
    - Lower final accuracy compared to more complex models.
    - Limited scalability.

12.2.2 Video recognition system:

First, I have tried LSTM model with background removal and after background removal mideapip to extract keypoints and train on them with LSTM model but I hade problem with the published datasets all of them was to small (each word have number of videos between 2 and 8) and that was to bad for training.

Second, I have tried to create my own dataset as all of the papers but there was problem when I split the video to frames some of the videos may have 40 frames and other may have 120 frames so if I fixed the length of the frames to 40 there will be data loss, so this solution did not work.

Third, I searched on dataset for another type of sign language and I found dataset for Brazilian sign language (BSL) and this dataset was already split to frames and I made VGG19 model to train on it.

VGG-19 Model (BSL)

- **Type:** VGG-19 (Convolutional Neural Network)
- **Dataset:** Brazilian Sign Language (BSL) Dataset
- **Training Time:** High
- **Hyperparameters:**
    - Learning Rate: 0.0001
    - Batch Size: 32
    - Epochs: 25
- **Performance Metrics:**
    - Accuracy: 98.3%
    - Precision: 98%
    - Recall: 98%
    - F1 Score: 98%
- **Strengths:**
    - High accuracy and robust performance.
    - Well-known architecture with proven performance in various tasks.
- **Weaknesses:**
    - Overfitting potential, as indicated by the significant drop in evaluation performance.
    - Requires significant computational resources.
    - Longer training time.

12.2.3 Video Request:

This part of the code was made to help people who did not know the sign language and want to learn it this part while help them.

This part work by select one word from the list of words that while appear to the user and the video while appear on the screen to this word so it will help people to know the sign language with easy way to increase the communication and make bride between deaf and mute people and normal people

13 Conclusions and Future Work

13.1 Summary

This project for developing system for translating sign language into text using various machine learning and deep learning models. The key achievements include:

1. **Model Development**: Three different models were implemented for American Sign Language (ASL), Arabic Sign Language (ARSL), and Brazilian Sign Language (BSL).
    - **CNN Model**: Showed good performance with an accuracy of 96.53% on the ASL dataset.
    - **VGG19 Model**: Achieved 100% training accuracy but exhibited overfitting with a significant drop in evaluation accuracy.
    - **EfficientNetB0 Model**: Provided the highest performance metrics with 99.7% accuracy, demonstrating its suitability for ASL recognition despite high computational requirements.
2. **Comparative Analysis**: A detailed comparison between SVM and EfficientNetB0 models for ASL recognition highlighted EfficientNetB0's superior accuracy but also its higher demand for computational resources.
3. **Video Recognition System**: Multiple trials were conducted using models like LSTM and VGG19 for video-based sign language recognition. The integration of background removal and feature extraction techniques improved model performance.
4. **Real-Time Application**: Developed a real-time sign language recognition system that uses webcam input to recognize and translate sign language gestures into text, enhancing user interaction and accessibility.

13.2 Future Work

Data Quality and Variety: The models' performance can be further enhanced by incorporating more diverse and high-quality datasets, especially for real-time applications where prediction accuracy can be variable.

Model Scalability: Exploring more scalable models and architectures that balance high accuracy with computational efficiency, making them feasible for deployment in resource-constrained environments.

Expanded Language Support: Extending the system to support additional sign languages and dialects, thereby broadening the scope and applicability of the technology.

User Interface and Experience: Enhancing the user interface to make it more intuitive and accessible, particularly for individuals with different levels of familiarity with technology.

Video Dataset Creation: Developing a comprehensive video dataset for sign language as current datasets are inadequate. This would significantly improve the performance and reliability of video-based sign language recognition models.

Real-Time Prediction Enhancement: Improving the real-time prediction capabilities to ensure more accurate and faster translations, making the system more effective for practical use.

14 References

[1] G Adewale, V. A., & Olamiti, D. A. (2018). *CONVERSION OF SIGN LANGUAGE TO TEXT AND SPEECH USING MACHINE LEARNING TECHNIQUES.* Edinburgh: Crawford University; University of Ibadan.

[2] Adeyanju, I., Bello, O., & Adegboye, M. (2021). Machine learning methods for signlanguage recognition: A critical review and analysis. Nigeria: Elsevier Ltd.

[3] Johnny, S., & Nirmala, S. J. (2021). Sign Language Translator Using Machine Learning. Singapore: Springer Nature Singapore.

[4] Mohammed, A., Amin, M., & Hefny, H. (2022). A Deep Learning Approach for Gloss Sign Language Translation using Transformer. *Computing and Communication* , 8.

[5] Reddy, G. S., & Dr.Parvathi.R. (2021). Sign Language to Text Conversion . *Advanced Trends in Computer Science and Engineering*, 4.

[6] Thakar, S., Shah, S., Shah, B., & Nimkar, A. V. (2022). *Sign Language to Text Conversionin Real Time.* Mumbai: Sardar Patel Institute of Technology.